

Natural Language Processing

Assignment 01 (Fall 2025)

Instructor: Ayesha Enayet

August 31, 2025

Question 1: Train BPE vs. WordPiece From Scratch & Compare

Learning Goals

- Understand how **Byte-Pair Encoding (BPE)** and **WordPiece** build subword vocabularies.
- Implement both tokenizers **from scratch** using only the Python standard library.
- Inspect the **first 10 merges** for each algorithm.
- Train both tokenizers on the q1_corpus_near1000 corpus available at https://drive.google.com/drive/folders/1po9naPNNrRhZl5rzvMd66MHu3svTerId?usp=drive_link.
- Compute token/subword **frequencies on the original corpus**, list the **top-10** for each tokenizer.
- Write a brief **comparative analysis** and **report challenges**.

Rules & Constraints

- Use only the Python standard library (`re`, `collections`, etc.). **No** external tokenization libraries (e.g., `tokenizers`, `sentencepiece`, `transformers`).
- **Vocabulary cap = 100** tokens for both methods (include single-character base tokens and any special markers).
- Print and clearly label:
 1. First **10 merges** for **BPE** and for **WordPiece**.
 2. **Top-10 tokens** (token, frequency) for each trained tokenizer on the **original corpus**.
- Follow the algorithmic variants discussed in class: BPE merges the most frequent *adjacent pair* (use an end-of-word marker, e.g., `</w>`); WordPiece uses greedy longest-match tokenization with a continuation prefix (e.g., `,`) and merges by your class scoring rule.

$$\text{score}(x, y) = \frac{\text{count}(x \triangleright y)}{\text{count}(x) \text{count}(y)}.$$

Mandatory Implementation Blocks (Must Appear Clearly in Your Code)

Your implementation must include separate functions for the following operations. These may call your helper functions.

1. **Corpus Reader:** implement `read_text_file(...)`. You may call helper functions.
2. **Single-Step Merge Implementations (BPE & WordPiece):** implement `bpe_merge_once(...)` and `wordpiece_merge_once(...)`; these must be explicitly invoked during training to print the first 10 merges. You may call helpers.
3. **Top-K Frequency Reporting:** implement `print_top_k(...)` to print the most frequent tokens/subwords for each trained tokenizer. You may call helpers.

What to Submit

1. **Code:** a single `bpe_wordpiece_from_scratch.py` containing the three required blocks above, BPE/WordPiece training & tokenization, and clearly printed outputs.
2. **Output File** (copy into your report as well):
 - First 10 merges for BPE and WordPiece (clearly labeled).
 - Top-10 tokens (token, frequency) for each trained tokenizer on the original corpus.
3. **Short report (PDF, 1–2 pages):**
 - Comparative analysis: early merges, morphological units, function words, affixes.
 - Challenges: tie-breaking, punctuation/Unicode, speed/memory, boundary effects, reproducibility.

Grading Rubric (100 pts)

- Correct BPE implementation (25 pts).
- Correct WordPiece implementation (25 pts).
- Report (Comparative analysis) (35 pts).
- Code quality & documentation (15 pts).

Academic Integrity: This is an individual assignment. Cite any resources you consult. Do not use existing tokenization libraries or copy implementations.

Question 2: Record Extraction Using Regex

Write a program that reads a text file containing 50 records in the form `{abc, 124, 23/45/67}` where each record contains exactly three fields: **name**, **phone**, and **date of birth (dob)**. The *order varies by record* (e.g., sometimes **name** first, sometimes **dob** first). Records are separated by curly braces (e.g., `{}`). Using **regular expressions**, extract the **name** and **dob** from each record and write a new file in which each line is **name,dob**.

Input format

- The file is UTF-8 text. Records appear like `{...}{...}{...}` and may span lines.
- Input file Question2_input URL: https://drive.google.com/drive/folders/1po9naPNNrRhZl5rzvMd66MHu3susp=drive_link
- Each record contains three comma-separated fields in *any order*:
 - **name**: alphabetic words (may include spaces, apostrophes, hyphens, dots).
 - **phone**: a number-like field (digits with optional separators, e.g., +92-300-1234567, 124).
 - **dob**: a date in *varying formats*, e.g., dd/mm/yyyy, d/m/yy, yyyy-mm-dd, dd-mm-yyyy, mm/dd/yyyy, 12 Mar 1980, March 12, 1980.

Tasks

1. **Read the file** and **extract records** using a regex that finds balanced curly-brace blocks.
2. **Within each record**, use **regex-based classification** to identify which token is `name`, which is `dob`, and which is `phone`.
 - Use a *date* alternation pattern that covers multiple formats (see “Regex sketch” below).
 - Use a *phone* pattern that primarily matches digit-heavy strings (allow + - () separators).
 - Treat the remaining token (that is not phone or date) as *name*; validate with a name regex.
3. **Normalize dates** to ISO YYYY-MM-DD. If the year is two digits, define a clear rule (e.g., 00-24 → 2000–2024; otherwise 1900–1999) and state it in a comment.
4. **Write output** to a new file (UTF-8): one record per line in the form `name,dob` (ISO). Preserve original name spacing/punctuation.

Constraints

- You **must** use Python’s `re` for field extraction/classification. No external parsing/tokenization libraries.
- You may use the Python standard library (e.g., `datetime`) for date normalization *after* regex extraction.
- Handle stray whitespace robustly (`\s*` around commas and inside braces). Ignore empty or malformed records gracefully.

Required functions (clearly separated in your code)

- `read_text(...)`
- `extract_records(...)`
- `classify_tokens(...)`
- `normalize_dob(...)`
- `write_pairs(...)`

Deliverables

- Code (with short comments).
- Output file.
- Report.

Report Submission (Regex Structuring Task)

What to include

1. Screenshot of Output File

Insert a clear screenshot of your final `name,dob` output file (one record per line).

2. Regular Expressions Used (with brief descriptions)

List every regex you used to (i) extract records, (ii) split/classify tokens, and (iii) detect dates/-phones/names. Provide a short, plain-English purpose for each.

3. Challenges Faced & Decisions Made

Briefly discuss:

- Ambiguities between phone vs. date vs. name and how your regex resolved them.
- Handling of whitespace, punctuation, two-digit years, and inconsistent date formats.
- Any rejected approaches (and why) and how you validated correctness.

Submission Format

- Submit a single PDF titled `Assign01_Report_YourName.pdf`.
- Include the screenshot (embedded), the regex table, and the challenges section.
- Keep it concise (1–2 pages). Ensure all text in screenshots is legible.

Grading (100 pts)

- Correct regex-based record extraction and token classification (50 pts)
- Accurate date normalization and output formatting (10 pts)
- Code organization and clarity (10 pts)
- Report (30 pts)

Post-Submission Demo (Required). After the submission deadline and the official announcement, each student **must** visit my office during office hours to demonstrate their code. Failure to complete the demo within the announced window will result in a penalty of 75 pts.