

* Notes :

- Query \rightarrow data retrieval.
- Transaction \rightarrow some data to be read/written in the database.
- Info stored by DBMS; Metadata
- Tuple \rightarrow sequence / list of values.
- atomic \rightarrow indivisible. / can't be broken down.
- Key \rightarrow Uniquely identify tuple
- SuperKey \rightarrow set of Keys.
 - Minimal SuperKey \rightarrow Candidate Key.
- Candidate Key: Several distinct attributes.
- Primary Key \rightarrow Candidate Key \rightarrow identifies tuples \rightarrow property of entire relation.
- Foreign Key \rightarrow link PK. to its relation

* Attributes :

Simple \rightarrow atomic

Composite Attributes \rightarrow divided into Subparts.

Multivalued: phone#, DOB

Derived: inferred / Calculated from another attribute \rightarrow age from DOB

★ Normalisation:

- DB free of anomalies

1NF: only atomic values / single values / set of values.

Resolution

- Eliminate repeating groups.

- Separate table for each set of related data.

2NF:

- Every non prime attribute A in R is fully functionally dependent on PK of R.

Test: (HS attributes = PK.

- If PK = Single attribute \rightarrow no need for test (provided there are no partial dependencies on PK).

PK1 PK2 NPA1 ✓ fully depends on PK.
└──────────┘

PK1 PK2 NPA2 X partially dependent on PK.

Resolution:

- NPA are associated only with the part of the PK they are dependent upon.
(fully functionality)

PK1 | PK2 | NPA1

PK1 | PK2 | NPA2

3NF: Transitive Dependency

$X \rightarrow Y \rightarrow Z$

hence

$X \rightarrow Z$

- No NPA is trans. dependent on


PK
 $X \rightarrow Y \rightarrow Z$

PK | NPA1 | NPA2

↑ ↑

PK | NPA1

NPA1 | NPA2



2NF \rightarrow PK contains multiple attributes \rightarrow no partial dep.

* Transaction Processing:

- The lost update problem.

\hookrightarrow an updated value gets lost as it is not written.

T_1	T_2
80 $r(x)$	
75 $X := X - N$	\rightarrow lost update.
	$r(x)$ 80
	$X := X + N$ 85

- The temporary update problem:

	T_1	T_2
TIME \downarrow	$r(x)$	
	$X = X - N$	
	$w(x)$	
		$r(x)$
		$X = X + M$
		$w(x)$
	abort.	

- Occurs when 1 transaction updates an item & fails \rightarrow but updated item is used by another transaction b4 item is changed or reverted back to its last value.

- Incorrect Summary%

- Transaction calculates some values b4 they are updated and others after they are updated.

example:

T_1	T_3
	Sum = 0
	$r(A)$
	$S = S + A ;$
$v(x)$	
$x = X - N$	
$w(x)$	
	$r(x) \rightarrow$ reads x after n is subtracted & y b4 n is added hence wrong summary
	$S = S + X$
	$r(y)$

$$S = S + \gamma$$

$r(Y)$

$Y = Y + N$

$w(Y)$

Unrepeatable read:

- T_1 reads the same item twice but item is changed by T_2 b/w 2 diff reads.

example \rightarrow Air line Seats.

* Transactions: \rightarrow Should either be

Atomicity \rightarrow performed in its entirety or not at all.
(Atomic Unit)

Consistency

Isolation

Durability.

$\rightarrow T_1$ should take DB from 1 consistent state to another if it is completely executed from beg. to end. w/o interfering w other transactions

Each T_n is unaware of other trans. executing simultaneously.

Transaction must operate without interference by any other T executing concurrently.

Changes to a System persist even if there are system failures.

Once a T completes successfully all its updates in the DB persist even if there is a system failure.

* Storage Structure:

Volatile \rightarrow does not survive system

crashes \rightarrow cache memory.

Non volatile \rightarrow survives system crashes \rightarrow secondary storage devices

Stable \rightarrow data is never lost.

Aborted \rightarrow T does not execute successfully.

\hookrightarrow if changes have been undone \rightarrow T has been rolled back

Committed \rightarrow T completes execution successfully
 \hookrightarrow effects can't be undone.

* Transaction states:

Active \rightarrow stays in this while executing.

Partially \rightarrow after final has been executed.

- \rightarrow order of execution.
- Serial Schedule: Instructions belonging to one transaction appear together in that schedule. \rightarrow always consistent
- Concurrent aren't always consistent.
if concurrent = serial \rightarrow serializable.

\rightarrow only on same var.
 $V-R \rightarrow$ no conflict
 $R-W, W-R, W-W \rightarrow$ yes conflict.

Swapping: If I and J (instructions of diff transactions) do not conflict, we can swap them to produce a new schedule S.

- If swapping can convert S into

serial \rightarrow Conflict equivalent.

View Serializability: S may not be Conflict serializable but it may produce the same view as a serial schedule.

* Concurrency Control :

— implemented through locks.

Shared \rightarrow read.

Exclusive \rightarrow read & write.

	S	X
S	T	F
X	F	F

Delayed Unlocking
 \hookrightarrow Consistent
Schedule.

Deadlock. \rightarrow locks are being requested w/o them being unlocked.
here \rightarrow deadlock.

Deadlocks > Inconsistency

— Unlocking too soon may give inconsistent
R.L. ...

states. (must in this basis)

- not unlocking before requesting =
deadlock.

↳ desirable as they can be handled by
rollback.

Starvation:

- exclusive locks may not be granted
if multiple transactions request

Shared locks (compatibility matrix)

→ Resolution: Sequential locking.

* 2 phase locking protocol:

- ensures serializability:

↳ growing phase → obtains locks → no
release.

Shrink → releases locks → wait
obtain, can't issue anymore lock requests.


- Deadlocks can still occur.

* Database backup:

1) Full backup:


involves copying entire data set. onto external disk.

- a lot of space, time, impractical daily.
- better weekly, biweekly.



2) Incremental:


- stores all files changed since last backup.
- last backup may be full, differential or incremental.
- resource friendly alternative to full.
- takes up less storage., faster




(full)

3) Differential backup:


- updates all data that has changed since the last full backup
- relies on most recent last full backup.
- faster than full
- takes up less space than full.




Mon




Tue



Wed



Thu



Fri

Summary:

full
copies entire
data

Tue:
full backup +
changes since
prev. backup

Diff:
full backup:
+ changes
since last
full backup

* Indexing:

- Many queries reference a small portion of records in a file here it is inefficient for system to read every tuple.

↳ works Just like an index in a text book.

★ Types:

Ordered: based on sorted ordering.

Hash: values are assigned according to hash function.

★ Parameters:

- Access types \rightarrow item search + range.

- Access time \rightarrow time taken to find particular data item

- Insertion time: insert a new data item (time taken to search + update).


- Deletion: time taken to delete + update

Ordered Index


Cluster: Sequentially ordered → primary Index.
Secondary: order different from Sequential → non cluster.

Index Sequential: Clustering Index on the search key.

Clustered:




Dense index: index entry appears for every search key value in the file.



Sparse: index entry appears for only some of the search-key values.
A relation is in sorted order

→ group of variation is in sorted order
i.e clustering index.




Dense is faster than sparse. (to search)


Sparse require less space.

↳ sparse index with 1 index entry per block.

Resolution:




Multilevel indices.



Outer


Inner.



2 level sparse.

Secondary:

much like dense: not stored sequentially.



Indexing \rightarrow efficiently retrieve records from db.

types: dense, sparse, clustered, non clustered.

If index is too big to fit in main memory, several level of indexing may be used.