

CS232L Operating Systems Lab

Lab 11: Introduction to POSIX Threads

CS Program
Habib University

Fall 2024

1 Introduction

A thread is a unit of execution. Each process has at least one thread. Sometimes multiple threads may be associated with a process.

In this lab you will learn:

1. a brief introduction to Posix Threads
2. create multiple threads
3. pass arguments to a thread
4. receive return values from a thread
5. wait for the termination of created threads

2 What is a Thread?

Threads are a mechanism that enable a process to perform multiple tasks concurrently. A single process can have multiple threads as shown in Figure 1 on page 2.

All of these threads are independently executing the same program, and they all share the same global memory, including the initialized data, uninitialized data, and heap segments. (A traditional UNIX process is simply a special case of a multithreaded processes; it is a process that contains just one thread.)

Threads in a process can execute concurrently. On a multi-processor system, the threads can execute in parallel¹.

In Linux the variable `errno` is maintained separately for each thread.

3 Thread creation

The function `pthread_create()`² creates a thread of the function passed to it as argument. The Listing 1 shows the use of `pthread_create` to execute the function `hello_fun()` in a separate thread. Note: Do not forget to link with the `pthread` library when compiling a program using POSIX threads by giving the option `-lpthread` to `gcc`.

¹While the same multiple processes can achieve the same effects, they have the following disadvantages:

- processes are expensive to create.
- it is more difficult to share data among processes.

This is why threads are sometimes referred to as light-weight processes.

²see manual page for the function

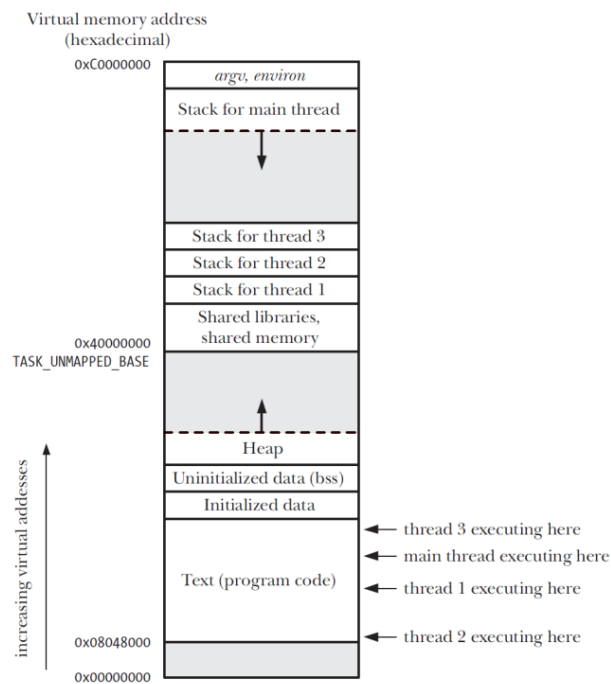


Figure 1: A Linux process running three threads.

```

1  /* hello_thread.c */
2  #include <stdio.h>
3  #include <stdlib.h>
4  #include <string.h>
5
6  #include <pthread.h>
7
8
9  void * hello_fun(void * args){
10
11     printf("Hello World!\n");
12
13     return NULL;
14 }
15
16 int main(int argc, char * argv[]) {
17
18     pthread_t t0, t1; //thread identifier
19
20     //create a new thread have it run the function hello_fun
21     pthread_create(&t0, NULL, hello_fun, NULL);
22     pthread_create(&t1, NULL, hello_fun, NULL);
23
24     //wait until the thread completes
25     pthread_join(t0, NULL);
26     pthread_join(t1, NULL);
27
28     return 0;
29 }

```

Listing 1: Creating a thread.

3.1 Exercise

- Create multiple threads in your program and make them perform different functions.

- Create multiple threads in your program and make them display their thread identifiers (TIDs) ³⁴ ⁵ and process identifiers (PIDs).

4 Passing arguments to threads

The Listing 2 shows how we can pass arguments to a thread when creating it.

```

1  /* hello_args_thread.c */
2  #include <stdio.h>
3  #include <stdlib.h>
4  #include <string.h>
5
6  #include <pthread.h>
7
8
9  void * hello_arg(void * args){
10
11     char * str = (char *) args;
12
13     printf("%s\n", str);
14
15     return NULL;
16 }
17
18 void * hello_arg_int(void * args){
19
20     int * str = (int *) args;
21
22     printf("%d\n", *str);
23
24     return NULL;
25 }
26
27 int main(int argc, char * argv[]) {
28
29     char hello [] = "Hello World!";
30     int x = 33;
31
32     pthread_t t1, t2; //thread identifier
33
34     //create a new thread that runs hello_arg with argument hello
35     pthread_create(&t1, NULL, hello_arg, hello);
36     pthread_create(&t2, NULL, hello_arg_int, &x);
37
38     //wait until the thread completes
39     pthread_join(t1, NULL);
40     pthread_join(t2, NULL);
41
42     return 0;
43 }

```

Listing 2: Passing arguments to threads.

4.1 Exercise

Make multiple threads and pass them different strings which they will each display along with their TIDs and PIDs.

³Checkout the function `pthread_self()` and `gettid()`

⁴You can google to see how you can print structures like `pthread_t`

⁵<https://stackoverflow.com/questions/34370172/the-thread-id-returned-by-pthread-self-is-not-the-something-as-the-kernel-thr>.

5 Waiting for threads

The `pthread_join()` function lets you wait for a thread to terminate. It is a blocking call which would block the calling thread until the thread specified as the join's argument has finished. Usually the model is that there is a main thread which creates all the other threads and then waits for all the threads to finish before exiting itself. However, nothing stops other threads from waiting for each other by calling `pthread_join()`.

If the main thread exits, all the threads created by it are joined even if they have not yet finished their execution. This is illustrated by the Listing 3. Observe the behaviour of this program by compiling and running it.

```
1  /* hello_pthread.bad.c */
2  #include <stdio.h>
3  #include <stdlib.h>
4  #include <string.h>
5
6  #include <pthread.h>
7
8
9  void * hello_fun () {
10
11     printf("Hello World!\n");
12
13     return NULL;
14 }
15
16 int main(int argc, char * argv []) {
17
18     pthread_t thread;
19
20     pthread_create(&thread, NULL, hello_fun, NULL);
21
22     //sleep(1);
23
24     return 0;
25 }
```

Listing 3: Failing to wait.

5.1 Exercise

- Correct the program in Listing 3 so that it displays the output.
- Create multiple threads in a program and see the effect of calling `pthread_exit()` and `exit()` in those threads.

6 Returning values from threads

Listing 4 shows how the second argument of `pthread_join()` can be used to receive the return value from a thread.

```
1  #include <stdio.h>
2  #include <stdlib.h>
3  #include <string.h>
4
5  #include <pthread.h>
6
7  static void *
8  threadFunc(void *arg)
9  {
10     char *s = (char *) arg;
11     printf("%s", s);
12     return (void *) strlen(s);
13 }
```

```

14
15 int
16 main(int argc, char *argv[])
17 {
18     pthread_t t1;
19     void *res;
20     int s;
21
22     s = pthread_create(&t1, NULL, threadFunc, "Hello world\n");
23     if (s != 0){
24         perror("error: pthread_create");
25         exit(EXIT_FAILURE);
26     }
27
28     printf("Message from main()\n");
29
30     s = pthread_join(t1, &res);
31     if (s != 0) {
32         perror("error: pthread_join");
33         exit(EXIT_FAILURE);
34     }
35
36     printf("Thread returned %ld\n", (long) res);
37     exit(EXIT_SUCCESS);
38 }

```

Listing 4: Returning values from threads.

6.1 Exercise

Listing 5 initializes a large array by writing 10 times a value in each of its indices.

```

1 #include <stdio.h>
2 #include <stdlib.h>
3 #include <string.h>
4 #include <time.h>
5
6
7 int
8 main(int argc, char *argv[])
9 {
10     //pthread_t t1;
11     void *res;
12     int s;
13
14     //srand(time(NULL));
15
16     //int a[1000000000];
17     int *a = malloc(1000000000*sizeof(int));
18     if (a == NULL) {
19         perror("error: memory failure");
20         exit(EXIT_FAILURE);
21     }
22
23
24     for(int i=0; i<1000000000; i++)
25         a[i%1000000000] = 1000000000;
26
27
28     free(a);
29     exit(EXIT_SUCCESS);
30 }

```

Listing 5: Initializing a large array (dummy.c).

- Compile and run the program to see how long it takes. See if you could make it run faster via threads.

- Create a global variable `counter`. Write a function that increments this variable 10 million times. Launch this function in 4 threads and wait in `main()` for these threads to finish and then display the value of `counter` in `main()`. Explain the observed value.

7 Other functions

Check out the documentation for functions and try to use them in your programs:

- `pthread_exit()`
- `pthread_self()`
- `pthread_equal()`
- `pthread_detach()`

References

- [1] Raymond Eric S. *"Basics of Unix Philosophy" The Art of Unix Programming*. Addison-Wesley, Professional.
- [2] BTYANT, O'HALLARON. *Computer Systems, A Programmer's Perspective*. Pearson.