

# CS232L Operating Systems Lab

## Lab 15: Signals in depth

CS Program  
Habib University

Fall 2024

### 1 Introduction

This lab gives you a lot more details about various different signals. For this lab, you are required to:

1. implement code of each example and task shared.
2. make a PDF file explaining the different uses of all the signals covered in this lab.

### 2 Generating Signals

Every signal has a symbolic name starting with **SIG**. The signal names are defined in **signal.h**, which must be included by any C program that uses signals. The names of the signals represent small integers greater than 0. Table 1 describes the required **POSIX** signals and lists their default actions. Two signals, **SIGUSR1** and **SIGUSR2**, are available for users and do not have a preassigned use. Some signals such as **SIGFPE** or **SIGSEGV** are generated when certain errors occur; other signals are generated by specific calls such as *alarm*.

Generate signals from the shell with the **kill** command. The name **kill** derives from the fact that, historically, many signals have the default action of terminating the process. The **signal\_name** parameter is a symbolic name for the signal formed by omitting the leading **SIG** from the corresponding symbolic signal name.

#### SYNOPSIS

```
kill -s signal_name pid ...  
kill -l [exit_status]
```

**Example 1** The following command is the traditional way to send signal number 9 (**SIGKILL**) to process 3423.

```
kill -9 3423
```

**Example 2** The following command sends the **SIGUSR1** signal to process 3423.

```
kill -s USR1 3423
```

**Example 3** The **kill -l** command gives a list of the available symbolic signal names. A system running Sun Solaris produced the following sample output.

```
% kill -l  
HUP INT QUIT ILL TRAP ABRT EMT FPE KILL BUS SEGV SYS PIPE ALRM TERM USR1 USR2 CLD PWR WINCH URG F
```

Table 1: The POSIX required signals.

signal	description	default action
SIGABRT	process abort	implementation dependent
SIGALRM	alarm clock	abnormal termination
SIGBUS	access undefined part of memory object	implementation dependent
SIGCHLD	child terminated, stopped or continued	ignore
SIGCONT	execution continued if stopped	continue
SIGFPE	error in arithmetic operation as in division by zero	implementation dependent
SIGHUP	hang-up (death) on controlling terminal (process)	abnormal termination
SIGILL	invalid hardware instruction	implementation dependent
SIGINT	interactive attention signal (usually Ctrl-C)	abnormal termination
SIGKILL	terminated (cannot be caught or ignored)	abnormal termination
SIGPIPE	write on a pipe with no readers	abnormal termination
SIGQUIT	interactive termination: core dump (usually Ctrl—)	implementation dependent
SIGSEGV	invalid memory reference	implementation dependent
SIGSTOP	execution stopped (cannot be caught or ignored)	stop
SIGTERM	termination	abnormal termination
SIGTSTP	terminal stop	stop
SIGTTIN	background process attempting to read	stop
SIGTTOU	background process attempting to write	stop
SIGURG	high bandwidth data available at a socket	ignore
SIGUSR1	user-defined signal 1	abnormal termination
SIGUSR2	user-defined signal 2	abnormal termination

## 2.1 Task

Call the `kill` function in a program to send a signal to a process. The `kill` function takes a process ID and a signal number as parameters. If the `pid` parameter is greater than zero, `kill` sends the signal to the process with that ID. If `pid` is 0, `kill` sends the signal to members of the caller's process group. If the `pid` parameter is -1, `kill` sends the signal to all processes for which it has permission to send. If the `pid` parameter is another negative value, `kill` sends the signal to the process group with group ID equal to `pid`.

### SYNOPSIS

```
#include <signal.h>
int kill(pid_t pid, int sig);
```

If successful, `kill` returns 0. If unsuccessful, `kill` returns -1 and sets `errno`. Table 2 lists the mandatory errors for `kill`.

Table 2: Mandatory errors of `kill` function.

error no.	cause
EINVAL	<code>sig</code> is an invalid or unsupported signal
EPERM	caller does not have the appropriate privileges
ESRCH	no process or process group corresponds to <code>pid</code>

A user may send a signal only to processes that he or she owns. For most signals, `kill` determines permissions by comparing the user IDs of caller and target. `SIGCONT` is an exception. For `SIGCONT`, user IDs are not checked if `kill` is sent to a process that is in the same session.

**Example 4** The following code segment sends the `SIGUSR1` signal to process 3423.

```
if ( kill(3243, SIGUSR1) == -1 )
    perror("Failed to send the SIGUSR1 signal");
```

Normally, programs do not hardcode specific process IDs such as 3243 in the `kill` function call. The usual way to find relevant process IDs is with `getpid`, `getppid`, `getgpid` or by saving the return value from `fork`.

**Example 5** This scenario sounds grim, but a child process can kill its parent by executing the following code segment.

```
if ( kill(getppid(), SIGTERM) == -1)
    perror ("Failed to kill parent");
```

A process can send a signal to itself with the `raise` function. The `raise` function takes just one parameter, a signal number.

#### SYNOPSIS

```
#include <signal.h>
int raise(int sig);
```

If successful, `raise` returns 0. If unsuccessful, `raise` returns a nonzero error value and sets `errno`. The `raise` function sets `errno` to `EINVAL` if `sig` is invalid.

**Example 6** The following statement causes a process to send the `SIGUSR1` signal to itself.

```
if (raise(SIGUSR1) != 0)
    perror("Failed to raise SIGUSR1");
```

A key press causes a hardware interrupt that is handled by the device driver for the keyboard. This device driver and its associated modules may perform buffering and editing of the keyboard input. Two special characters, the `INTR` and `QUIT` characters, cause the device driver to send a signal to the foreground process group. A user can send the `SIGINT` signal to the foreground process group by entering the `INTR` character. This user settable character is often `Ctrl-C`. The user settable `QUIT` character sends the `SIGQUIT` signal.

**Example 7** The `stty -a` command reports on the characteristics of the device associated with standard input, including the settings of the signal-generating characters.

The terminal in Example 7 interprets `Ctrl-C` as the `INTR` character. The `QUIT` character (`Ctrl—` above) generates `SIGQUIT`. The `SUSP` character (`Ctrl-Z` above) generates `SIGSTOP`, and the `DSUSP` character (`Ctrl-Y` above) generates `SIGCONT`.

**Example 8 Code: `simplealarm.c`** Since the default action for `SIGALRM` is to terminate the process, the following program runs for approximately ten seconds of wall-clock time.

```
#include <unistd.h>
int main(void) {
    alarm(10);
    for ( ; ; ) ;
}
```

## 3 Blocking and Unblocking a Signal

A process can temporarily prevent a signal from being delivered by blocking it. Blocked signals do not affect the behavior of the process until they are delivered. The process *signal mask* gives the set of signals that are currently blocked. The signal mask is of type `sigset_t`.

Blocking a signal is different from ignoring a signal. When a process blocks a signal, the operating system does not deliver the signal until the process unblocks the signal. A process

blocks a signal by modifying its signal mask with `sigprocmask`. When a process ignores a signal, the signal is delivered and the process handles it by throwing it away. The process sets a signal to be ignored by calling `sigaction` with a handler of `SIG_IGN`.

Signal sets are manipulated by the five functions listed in the following synopsis box. The first parameter for each function is a pointer to a `sigset_t`. The `sigaddset` adds `signo` to the signal set, and the `sigdelset` removes `signo` from the signal set. The `sigemptyset` function initializes a `sigset_t` to contain no signals; `sigfillset` initializes a `sigset_t` to contain all signals. Initialize a signal set by calling either `sigemptyset` or `sigfillset` before using it. The `sigismember` reports whether `signo` is in a `sigset_t`.

#### SYNOPSIS

```
#include <signal.h>
int sigaddset(sigset_t *set, int signo);
int sigdelset(sigset_t *set, int signo);
int sigemptyset(sigset_t *set);
int sigfillset(sigset_t *set);
int sigismember(const sigset_t *set, int signo);
```

The `sigismember` function returns 1 if `signo` is in `*set` and 0 if `signo` is not in `*set`. If successful, the other functions return 0. If unsuccessful, these other functions return `-1` and set `errno`.

**Example 9** The following code segment initializes signal set `twosigs` to contain exactly the two signals `SIGINT` and `SIGQUIT`.

```
if ((sigemptyset(&twosigs) == -1) || (sigaddset(&twosigs, SIGINT) == -1) ||
    (sigaddset(&twosigs, SIGQUIT) == -1))
    perror("Failed to set up signal mask");
```

A process can examine or modify its process signal mask with the `sigprocmask` function. The `how` parameter is an integer specifying the manner in which the signal mask is to be modified. The `set` parameter is a pointer to a signal set to be used in the modification. If `set` is `NULL`, no modification is made. If `oset` is not `NULL`, the `sigprocmask` returns in `*oset` the signal set before the modification.

#### SYNOPSIS

```
#include <signal.h>
int sigprocmask(int how, const sigset_t *restrict set, sigset_t *restrict oset);
```

If successful, `sigprocmask` returns 0. If unsuccessful, `sigprocmask` returns `-1` and sets `errno`. The `sigprocmask` function sets `errno` to `EINVAL` if `how` is invalid. The `sigprocmask` function should only be used by a process with a single thread. When multiple threads exist, the `pthread_sigmask` function should be used.

The `how` parameter, which specifies the manner in which the signal mask is to be modified, can take on one of the following three values.

**SIG\_BLOCK**: add a collection of signals to those currently blocked

**SIG\_UNBLOCK**: delete a collection of signals from those currently blocked

**SIG\_SETMASK**: set the collection of signals being blocked to the specified set

Keep in mind that some signals, such as `SIGSTOP` and `SIGKILL`, cannot be blocked. If an attempt is made to block these signals, the system ignores the request without reporting an error.

**Example 10** The following code segment adds `SIGINT` to the set of signals that the process has blocked.

```
sigset_t newsigset;
if ((sigemptyset(&newsigset) == -1) || (sigaddset(&newsigset, SIGINT) == -1))
    perror("Failed to initialize the signal set");
```

```

else if (sigprocmask(SIG_BLOCK, &newsigset, NULL) == -1)
    perror("Failed to block SIGINT");

```

If `SIGINT` is already blocked, the call to `sigprocmask` has no effect.

### 3.1 Task

Write a program that displays a message, blocks the `SIGINT` signal while doing some work, unblocks the signal, and does more work. The program repeats this sequence continually in a loop.

If a user enters `Ctrl-C` while `SIGINT` is blocked, program finishes the calculation and prints a message before terminating. If a user types `Ctrl-C` while `SIGINT` is unblocked, the program terminates immediately.

The function `makepair` of code given below takes two pathnames as parameters and creates two named pipes with these names. If successful, `makepair` returns 0. If unsuccessful, `makepair` returns `-1` and sets `errno`. The function blocks all signals during the creation of the two pipes to be sure that it can deallocate both pipes if there is an error. The function restores the original signal mask before the return. The if statement relies on the conditional left-to-right evaluation of `&&` and `||`.

### 3.2 Task

Is it possible that after a call to `makepair`, `pipe1` exists but `pipe2` does not?

**Code: `makepair.c`** A function that blocks signals while creating two pipes.

```

#include <errno.h>
#include <signal.h>
#include <unistd.h>
#include <sys/stat.h>

#define R_MODE (S_IRUSR | S_IRGRP | S_IROTH)
#define W_MODE (S_IWUSR | S_IWGRP | S_IWOTH)
#define RW_MODE (R_MODE | W_MODE)

int makepair(char *pipe1, char *pipe2)
{
    sigset_t blockmask;
    sigset_t oldmask;
    int returncode = 0;

    if (sigfillset(&blockmask) == -1)
        return -1;
    if (sigprocmask(SIG_SETMASK, &blockmask, &oldmask) == -1)
        return -1;
    if (((mkfifo(pipe1, RW_MODE) == -1) && (errno != EEXIST)) ||
        ((mkfifo(pipe2, RW_MODE) == -1) && (errno != EEXIST)))
    {
        returncode = errno;
        unlink(pipe1);
        unlink(pipe2);
    }

    if ((sigprocmask(SIG_SETMASK, &oldmask, NULL) == -1) && !returncode)
        returncode = errno;
    if (returncode)

```

```

    {
        errno = returncode;
        return -1;
    }
    return 0;
}

```

### 3.3 Task

1. Does a `makepair` return value of 0 guarantee that FIFOs corresponding to `pipe1` and `pipe2` are available on return?
2. Write a program in which the parent blocks all signals before forking a child process to execute an `ls` command.

**Code: password.c** A function that retrieves a user password.

```

#include <errno.h>
#include <fcntl.h>
#include <signal.h>
#include <stdio.h>
#include <string.h>
#include <termios.h>
#include <unistd.h>
#include "restart.h"

int setecho(int fd, int onflag);
int password(const char *prompt, char *passbuf, int passmax)
{
    int fd;
    int firsterrno = 0;
    sigset_t signew, sigold;
    char termbuf[L_ctermid];

    if (ctermid(termbuf) == NULL)
    {
        /* find the terminal name */ errno = ENODEV;
        return -1;
    }

    if ((fd = open(termbuf, O_RDONLY)) == -1)
        /* open descriptor to terminal */ return -1;
    if ((sigemptyset(&signew) == -1) || /* block SIGINT, SIGQUIT and SIGTSTP */ (sigaddset(&signew, SIGINT) == -1) || (sigaddset(&signew, SIGQUIT) == -1) || (sigaddset(&signew, SIGTSTP) == -1) || (sigprocmask(SIG_SETMASK, &signew, &sigold) == -1) || (setecho(fd, 0) == -1))
    {
        /* set terminal echo off */ firsterrno = errno;
        sigprocmask(SIG_SETMASK, &sigold, NULL);
        r_close(fd);
        errno = firsterrno;
        return -1;
    }

    if ((r_write(STDOUT_FILENO, (char *)prompt, strlen(prompt)) == -1) || (getline(fd, passbuf, NULL) == -1))
        /* read password */
        firsterrno = errno;
}

```

```

else
    passbuf[strlen(passbuf) - 1] = 0; /* remove newline */
if ((setecho(fd, 1) == -1) && !firsterrno) /* turn echo back on */
    firsterrno = errno;
if ((sigprocmask(SIG_SETMASK, &sigold, NULL) == -1) && !firsterrno)
    firsterrno = errno;
if ((r_close(fd) == -1) && !firsterrno)
    /* close descriptor to terminal
    */
    firsterrno = errno;
return firsterrno ? errno = firsterrno, -1: 0;
}

```

The `password` function blocks `SIGINT`, `SIGQUIT` and `SIGTSTP` while terminal echo is set off, preventing the terminal from being placed in an unusable state if one of these signals is delivered to the process while this function is executing.

## 4 Catching and Ignoring Signals - `sigaction`

The `sigaction` function allows the caller to examine or specify the action associated with a specific signal. The `sig` parameter of `sigaction` specifies the signal number for the action. The `act` parameter is a pointer to a `struct sigaction` structure that specifies the action to be taken. The `oact` parameter is a pointer to a `struct sigaction` structure that receives the previous action associated with the signal. If `act` is `NULL`, the call to `sigaction` does not change the action associated with the signal. If `oact` is `NULL`, the call to `sigaction` does not return the previous action associated with the signal.

### SYNOPSIS

```

#include <signal.h>
int sigaction(int sig, const struct sigaction *restrict act, struct sigaction *restrict oact);

```

If successful, `sigaction` returns 0. If unsuccessful, `sigaction` returns -1 and sets `errno`. Table 3 lists the mandatory errors for `sigaction`. The `struct sigaction` structure must have at least

Table 3: Mandatory errors of `sigaction`.

error no.	cause
EINVAL	<code>sig</code> is an invalid signal number, or attempt to catch a signal that cannot be caught, or attempt to ignore a signal that cannot be ignored
ENOTSUP	<code>SA_SIGINFO</code> bit of <code>sig_flags</code> is set and the implementation does not support POSIX:RTS or POSIX:X

the following members.

```

1  struct sigaction
2  {
3      void (*sa_handler)(int); /* SIG_DFL, SIG_IGN or pointer to function */
4      sigset_t sa_mask; /* additional signals to be blocked during execution of
5                          handler */
6      int sa_flags; /* special flags and options */
7      void (*sa_sigaction)(int, siginfo_t *, void *); /* realtime handler */
8  };

```

Listing 1: `sigaction` struct members.

The storage for `sa_handler` and `sa_sigaction` may overlap, and an application should use only one of these members to specify the action. If the `SA_SIGINFO` flag of the `sa_flags` field is cleared, the `sa_handler` specifies the action to be taken for the specified signal.

**Example 11** The following code segment sets the signal handler for `SIGINT` to `mysighand`.

```

struct sigaction newact;
newact.sa_handler = mysighand; /* set the new handler */
newact.sa_flags = 0; /* no special options */
if ((sigemptyset(&newact.sa_mask) == -1) || (sigaction(SIGINT, &newact, NULL) == -1))
    perror("Failed to install SIGINT signal handler");

```

A signal handler is an ordinary function that returns `void` and has one integer parameter. When the operating system delivers the signal, it sets this parameter to the number of the signal that was delivered. Most signal handlers ignore this value, but it is possible to have a single signal handler for many signals. The usefulness of signal handlers is limited by the inability to pass values to them.

Two special values of the `sa_handler` member of `struct sigaction` are `SIG_DFL` and `SIG_IGN`. The `SIG_DFL` value specifies that `sigaction` should restore the default action for the signal. The `SIG_IGN` value specifies that the process should handle the signal by ignoring it (throwing it away).

**Example 12** The following code segment causes the process to ignore `SIGINT` if the default action is in effect for this signal.

```

struct sigaction act;
if (sigaction(SIGINT, NULL, &act) == -1) /* Find current SIGINT handler */
    perror("Failed to get old handler for SIGINT");
else if (act.sa_handler == SIG_DFL)
{
    /* if SIGINT handler is default */
    act.sa_handler = SIG_IGN; /* set new SIGINT handler to ignore */
    if (sigaction(SIGINT, &act, NULL) == -1)
        perror("Failed to ignore SIGINT");
}

```

**Example 13** The following code segment sets up a signal handler that catches the `SIGINT` signal generated by `Ctrl-C`.

```

void catchctrlc(int signo)
{
    char handmsg[] = "I found Ctrl-C\n";
    int msglen = sizeof(handmsg);
    write(STDERR_FILENO, handmsg, msglen);
}
...
struct sigaction act;
act.sa_handler = catchctrlc;
act.sa_flags = 0;
if ((sigemptyset(&act.sa_mask) == -1) || (sigaction(SIGINT, &act, NULL) == -1))
    perror("Failed to set SIGINT to handle Ctrl-C");

```

## 4.1 Task

Why didn't Example 12 use `fprintf` or `strlen` in the signal handler?

**Example 14** The following code segment sets the action of `SIGINT` signal to the default.

```

struct sigaction newact;
newact.sa_handler = SIG_DFL; /* new handler set to default */
newact.sa_flags = 0; /* no special options */
if ((sigemptyset(&newact.sa_mask) == -1) || (sigaction(SIGINT, &newact, NULL) == -1))
    perror("Failed to set SIGINT to the default action");

```



**Example 15: testignored.c** The following function takes a signal number parameter and returns 1 if that signal is ignored and 0 otherwise.

```
#include <signal.h>
#include <stdio.h>

int testignored(int signo)
{
    struct sigaction act;
    if ((sigaction(signo, NULL, &act) == -1) || (act.sa_handler != SIG_IGN))
        return 0;
    return 1;
}
```

## References

- [1] M. Kerrisk. *The Linux Programming Interface*. No Starch Press Series. No Starch Press, 2010. isbn: 9781593272203. url: <https://books.google.com.pk/books?id=2SAQAQAAQBAJ>.
- [2] W.R. Stevens and S.A. Rago. *Advanced Programming in the UNIX Environment: Advanced Programming in the UNIX Environment*. Addison-Wesley Professional Computing Series. Pearson Education, 2013. isbn: 9780321638007. url: <https://books.google.gl/books?id=kCTMFpEcIOWC>.