

- ① O_CREAT
- ② O_RDONLY
- ③ O_WRONLY

Date: _____

- int fd = open(): system call returns file descriptor (fd)
- file descriptor is an int, may be considered as a 'capability', allowing one to access files.

```

prompt > echo hello > foo
prompt > cat foo
hello
prompt >
    
```

an example of
reading & writing
foo.

How does the cat program access the file foo?

We can use strace to trace the system calls made by program.

prompt > strace cat foo

...

open ("foo", O_RDONLY | O_LARGEFILE)

= 3

file descriptor.

read (3, "hello\n", 4096)

= 6

length of string.

write (1, "hello\n", 6) → size of buffer:

= 6

standard output.

hello

read (3, "", 4096)

= 0 || no bytes left in file.

close (3)

= 0

...

prompt >

* file descriptor helps identify the file uniquely within same process.

Date: _____

- each running process already has three files open
 - (i) standard input [0]
 - (ii) standard output [1] and (iii)
 - standard error [2].
- "cat" stands for concatenate.
 - `read(file desc., buffer pointer, size of buffer)` the value it returns is length of string in above example "hello\n" is of 6 length.
 - `write(file desc., buffer pointer, size of buffer)` has file desc. 1 for standard output and it also returns value 6 the size of string.
 - `close(3)` closes the file with file desc 3.

What if- we want to read from random offset within a document?

? `off_t lseek(int fd, off_t offset, int whence);`

 ↓ ↓ ↓
 file offset whence
 descriptor in bytes → determine how
 seek is performed
 → from start
 (SEEK_SET)

fork() → from current
 (SEEK_CUR)

→ here a parent process creates a
child process with `fork()` → from end
 (SEEK_END)

→ the parent then waits for child to
complete.

→ finally, the parent continues from where
the child completed its work.

(58-63)

Date: _____

(Refer to slide 57)

`int fd=open("file.txt", O_RDONLY)` → opens file in read only mode, returns fd if successfully opened.

`assert(fd>=0)` → if file doesn't open, program would terminate with assertion failure.

`int rc=fork()` → to create new process/child process returns 0 & returns PID of child process in parent process.

`if (rc==0){ ... }` → creation of child process.

`else if (rc>0) { ... }` → parent process, the `wait(NULL)` function called making parent process to wait for child function.

Date: _____

RENAMING FILES

rename (char * old, char * new)

- rename a file to a different name.
- it implemented as an atomic call.

prompt > mv foo bar. // change from foo to bar.

flushes all modified

file data & metadata

from buffers in int fint fd = open ("foo.txt", O_WRONLY | O_CREATE | O_TRUNC);

RAM to physical

disk write (fd, buffer, size);

fsync (fd);

close (fd);

rename ("foo.txt", "foo.txt.tmp")

(the order of flags
does not matter.)

- to get information about files-

stat(), **fstat()** - show file metadata.

→ The stat information,
is displayed in a 'inode'

is information
about each file.

structure.

REMOVING FILES

rm is Linux command to remove file.

- rm call **unlink()** to remove file.

Date: _____

MAKING DIRECTORIES

mkdir() - make directory

- when a directory is created, it is empty
- empty directory have two entries . (itself), .. (parent)

prompt> ls -a
· / .. /

DELETING DIRECTORIES

rmdir() - delete directory

- requires that directory be empty
- if you call rmdir() to a non-empty, it will fail

HARD LINKS

link(old pathname , new one) - link new file to old.

- create another way to refer to same file.
- the command-line link program: ln
- the way link works:
 - create another name in directory
 - refer it to same inode number of original file.
 - the file is not copied anyway
- now we just have two human names (file and file2) that both refer to same file.

* a hardlink is an additional name for an existing file on same filesystem.
→ direct reference to file's inode.

* symbolic link is a special type of file that serves as a reference or pointer to another file/directory.
→ has their own inodes and store path to the target file.

Date: _____

→ The result of link()

→ two files have same inode number, but two human name (file, file2)

→ there is no difference b/w file and file2.

→ Thus, to remove a file, we call unlink()

→ when unlink() is called reference count decrements

→ if reference count reaches zero, the filesystem free the inode and related data ~~blocks~~ → truly "delete" the file.

SYMBOLIC LINKS

→ Symbolic Links is more useful than Hard Link.

→ Hard Link cannot create to directory

→ Hard Link cannot create to file to other partition.

What is the difference b/w Symbolic Link and Hard Link?

→ Symbolic links are a third type file system knows about.

→ The size of symbolic links (file2) is 4 bytes.

→ A symbolic link holds the pathname of linked-to file as data of link file.

→ if we link to larger pathname, our link file would be bigger.

* Hard links remains valid if original file is deleted; file data not removed until all links are deleted. Symbolic links become 'dangling' if target file is moved/deleted as they store path to file.

M

Date: _____

→ Dangling reference.

→ When remove original file, symbolic link points nothing.

MAKING AND MOUNTING FILE SYSTEM

mkfs tool : make file system.

→ Write an empty file system starting with a root directory, on to disk partition.

→ Input:

→ a device (such as disk partition e.g /dev/sda1)

→ a file system type (e.g ext3)

mount()

→ take an existing directory as a target mount point

→ essentially paste a new file system auto directory tree at point

CHAPTER 40 : File System Implementation

→ There are two different aspects to implement filesystem.

① Data Structures.

* What types of on-disk structures are utilized by the file system to organize its data & metadata?

② Access Methods

* How does it map the calls made by process as open(), read() and write() etc.

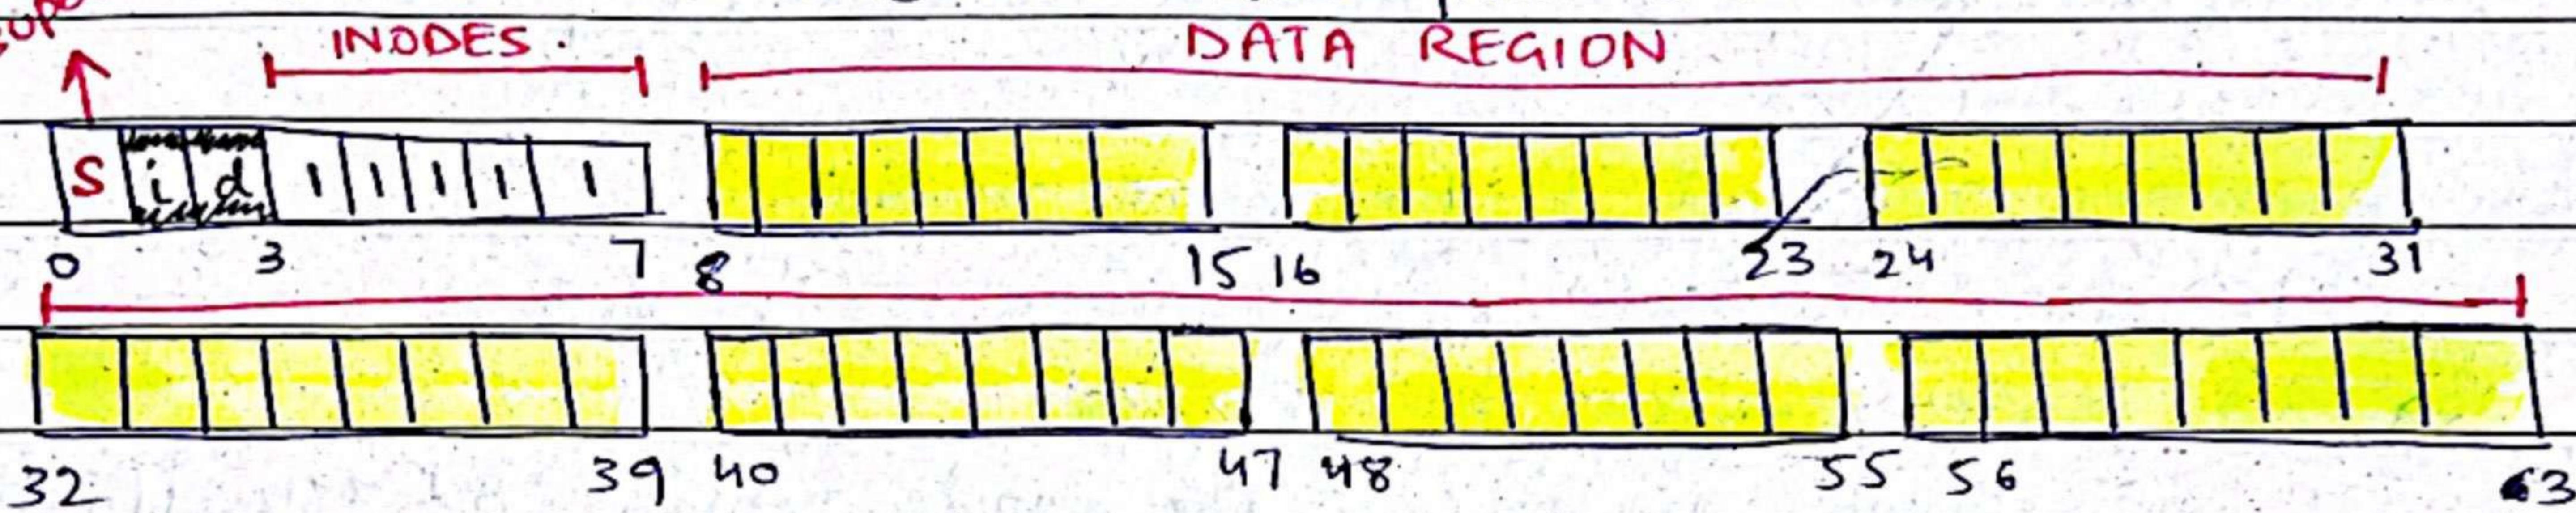
Date: _____

* Which structures are read during the execution of particular system call?

→ Divide the disk into blocks.

→ block size is 4KB.

→ the blocks are addressed from 0 to N-1.



→ filesystem has to track which data block comprise a file, the size of file, its owner etc.

How we store these inodes in file system? stores information about file/directory, contains metadata about files/directories

→ reserve some space for **inode table**.

→ this holds an array of on-disk inodes.

→ Ex) inode tables : 3 ~ 7, inode size : 256 bytes.

→ 4-KB block can hold 16 inodes.

→ the filesystem contains 80 inodes (maximum number of files)

→ The Bitmap is to track whether nodes or data blocks are free, or allocated.

Date: _____

- each **bitmap** bit indicates free (0) or in-use (1)
 - data bitmap: for data region
 - inode bitmap: for inode table. (^{tells which nodes are} free and which are allocated)
this
- **Superblock** contains information for particular file system
 - Ex) the number of inodes, begin location of inode table etc.
- Each inode is referred to by inode number.
 - filesystem calculate where inode is on disk.
 - Ex) inode number: 32.

① calculate offset into the inode region ($32 \times \text{sizeof(inode)}$)
(256 bytes) = 8192.

② Add start address of inode table (12KB) +
inode region (8 KB) = 20KB.

blk : $(\text{inumber} * \text{sizeof(inode)}) / \text{blocksize}$

sector : $((\text{blk} * \text{blocksize}) + \text{inodeStart Addr}) / \text{sectorsize}$.

inode have all of the information about a file.

- file type (regular file, directory, etc)
- size, the number of blocks allocated to it
- protection information (who owns the file, who can access, etc).
- time information.

Date: _____

- To support bigger files, we use multi-level index.
- Indirect pointer points to a block that contains more pointers
 - mode have fixed number of direct pointers (12) and single indirect pointer:
 - if a file grows large enough, an indirect block is allocated, inode's slot for an indirect pointer is set to point to it
 - $(12 + 1024) \times 4K$ or 4144KB.
- Double indirect pointer. points to a block that contains indirect blocks
 - Allow file to grow with an additional 1024×1024 or 1 million 4KB blocks.
- Triple indirect pointer. points to block that contains double indirect blocks
 - Multi-level index approach to pointing to file blocks.
 - Ex) twelve direct pointers, a single and a double indirect block.
 - over 4GB in size $(12 + 1024 + 1024^2) \times 4KB$.
- Many file system use multi-level index
 - Linux EXT2, EXT3, NetApp's WAFL, Unix file system.
 - Linux EXT4 use extents instead of simple pointers.

Date: _____

DIRECTORY ORGANIZATION

- directory contains a list of (entry name, mode number) pairs
- each directory has two extra files ".dot" for current directory and ".. dot dot" for parent directory.
- for example, dir has three files (foo, bar, foobar)

inum	reclen	strlen	name
5	4	2	.
2	4	3	..
12	4	4	foo
13	4	4	bar
24	8	7	foobar

- filesystem tracks which inode & data block are free or not
- in order to manage free space, we have two simple bitmaps.
- When the file is newly created, it allocates inode by searching the inode bitmap and update on-disk bitmap.
- Pre-allocation policy is commonly used for allocate contiguous blocks.

Date: _____

READING FILE FROM DISK

`open (" /foo/bar ", O_RDONLY)`

- traverse the pathname & thus locate desired inode.
- begin at the root of filesystem (/)
- in most Unix file systems, the root inode number is 2. *important.
- filesystem reads in block that contains inode number 2.
- look inside of it to find pointer to data blocks,
~~# with contents of root~~.
- by reading in one or more directory data blocks
it will find "foo" directory.
- traverse recursively the path name until the desired
inode ("bar")
- check file permissions, allocate a file descriptor
for this process & returns file descriptor to user.

- issue `read()` to read from the file.
- read in the first block of file, consulting the inode
to find location of such block.
- update the inode with a new last accessed
time.
- update in-memory open file table for file
descriptor, the file offset.

Date: _____

→ When the file is closed

→ file descriptor should be deallocated, but for now, that is all file system really needs to do.
No disk I/O take place.

→ Issue write() to update file with new contents

→ File may allocate a block (unless the block is being written)

→ need to update data block, data bitmap.

→ generates five I/Os.

① to read data bitmap

② to write bitmap

③ two more to read & write inode

④ are to write actual block itself.

CACHING & BUFFERING

→ Reading & writing files are EXPENSIVE, incurring many I/Os.

→ for e.g. long pathname

→ literally perform hundreds of reads just to open file.

→ In order to reduce I/O traffic, file systems aggressively use system memory (DRAM) to cache.

→ modern system use dynamic partitioning approach, unified page cache.

→ Read I/O can be avoided by large cache.

Date: _____

- Write traffic has to go to disk for persistent, thus, cache does not reduce write I/Os.
- filesystem use write buffering for write performance benefits
 - delaying writes
 - by buffering number of writes in memory, filesystem can then schedule the subsequent I/Os.
 - by avoiding writes.
- Some application force flush data to disk by calling `fsync()` or direct I/O.

CHAPTER 4: The Abstraction: The Process.

- CPU Virtualizing
 - The OS can promote the illusion that many virtual CPUs exist
 - Time Sharing: Running one process, then stopping it and running another.
 - The potential cost is **performance**.

PROCESS

A process is a **running program**.

- comprising of a process:
 - **Memory** instructions
 - data section
- Registers
 - program counter.
 - stack pointer.

Date: _____

Process API

→ These APIs are available on any modern OS.

→ Create

→ Create a new process to run a program.

→ Destroy

→ Halt a runaway process

→ Wait

→ Wait for process to stop running

→ Miscellaneous Control.

→ Some kind of method to suspend a process and then resume it.

→ Status

→ Get some status info about process

PROCESS CREATION

① Load a program code into memory into the address space of process.

→ programs initially reside on disk in executable format

→ OS perform loading process **lazily**.

→ loading pieces of code or data only as they are needed during program execution.

* The code/text segment is a region of process's virtual address space reserved for holding executable code.

Date: _____

② The program's run-time stack is allocated.

→ Use the stack for local variables, function parameters and return address.

→ Initialize the stack with arguments → argc and argv array of main() function.

③ The program's heap is created.

→ Used for explicitly requested dynamically allocated data.

→ Program request such space by calling malloc() and free it by calling free().

④ The OS do some other initialization tasks.

→ Input/output (I/O) setup.

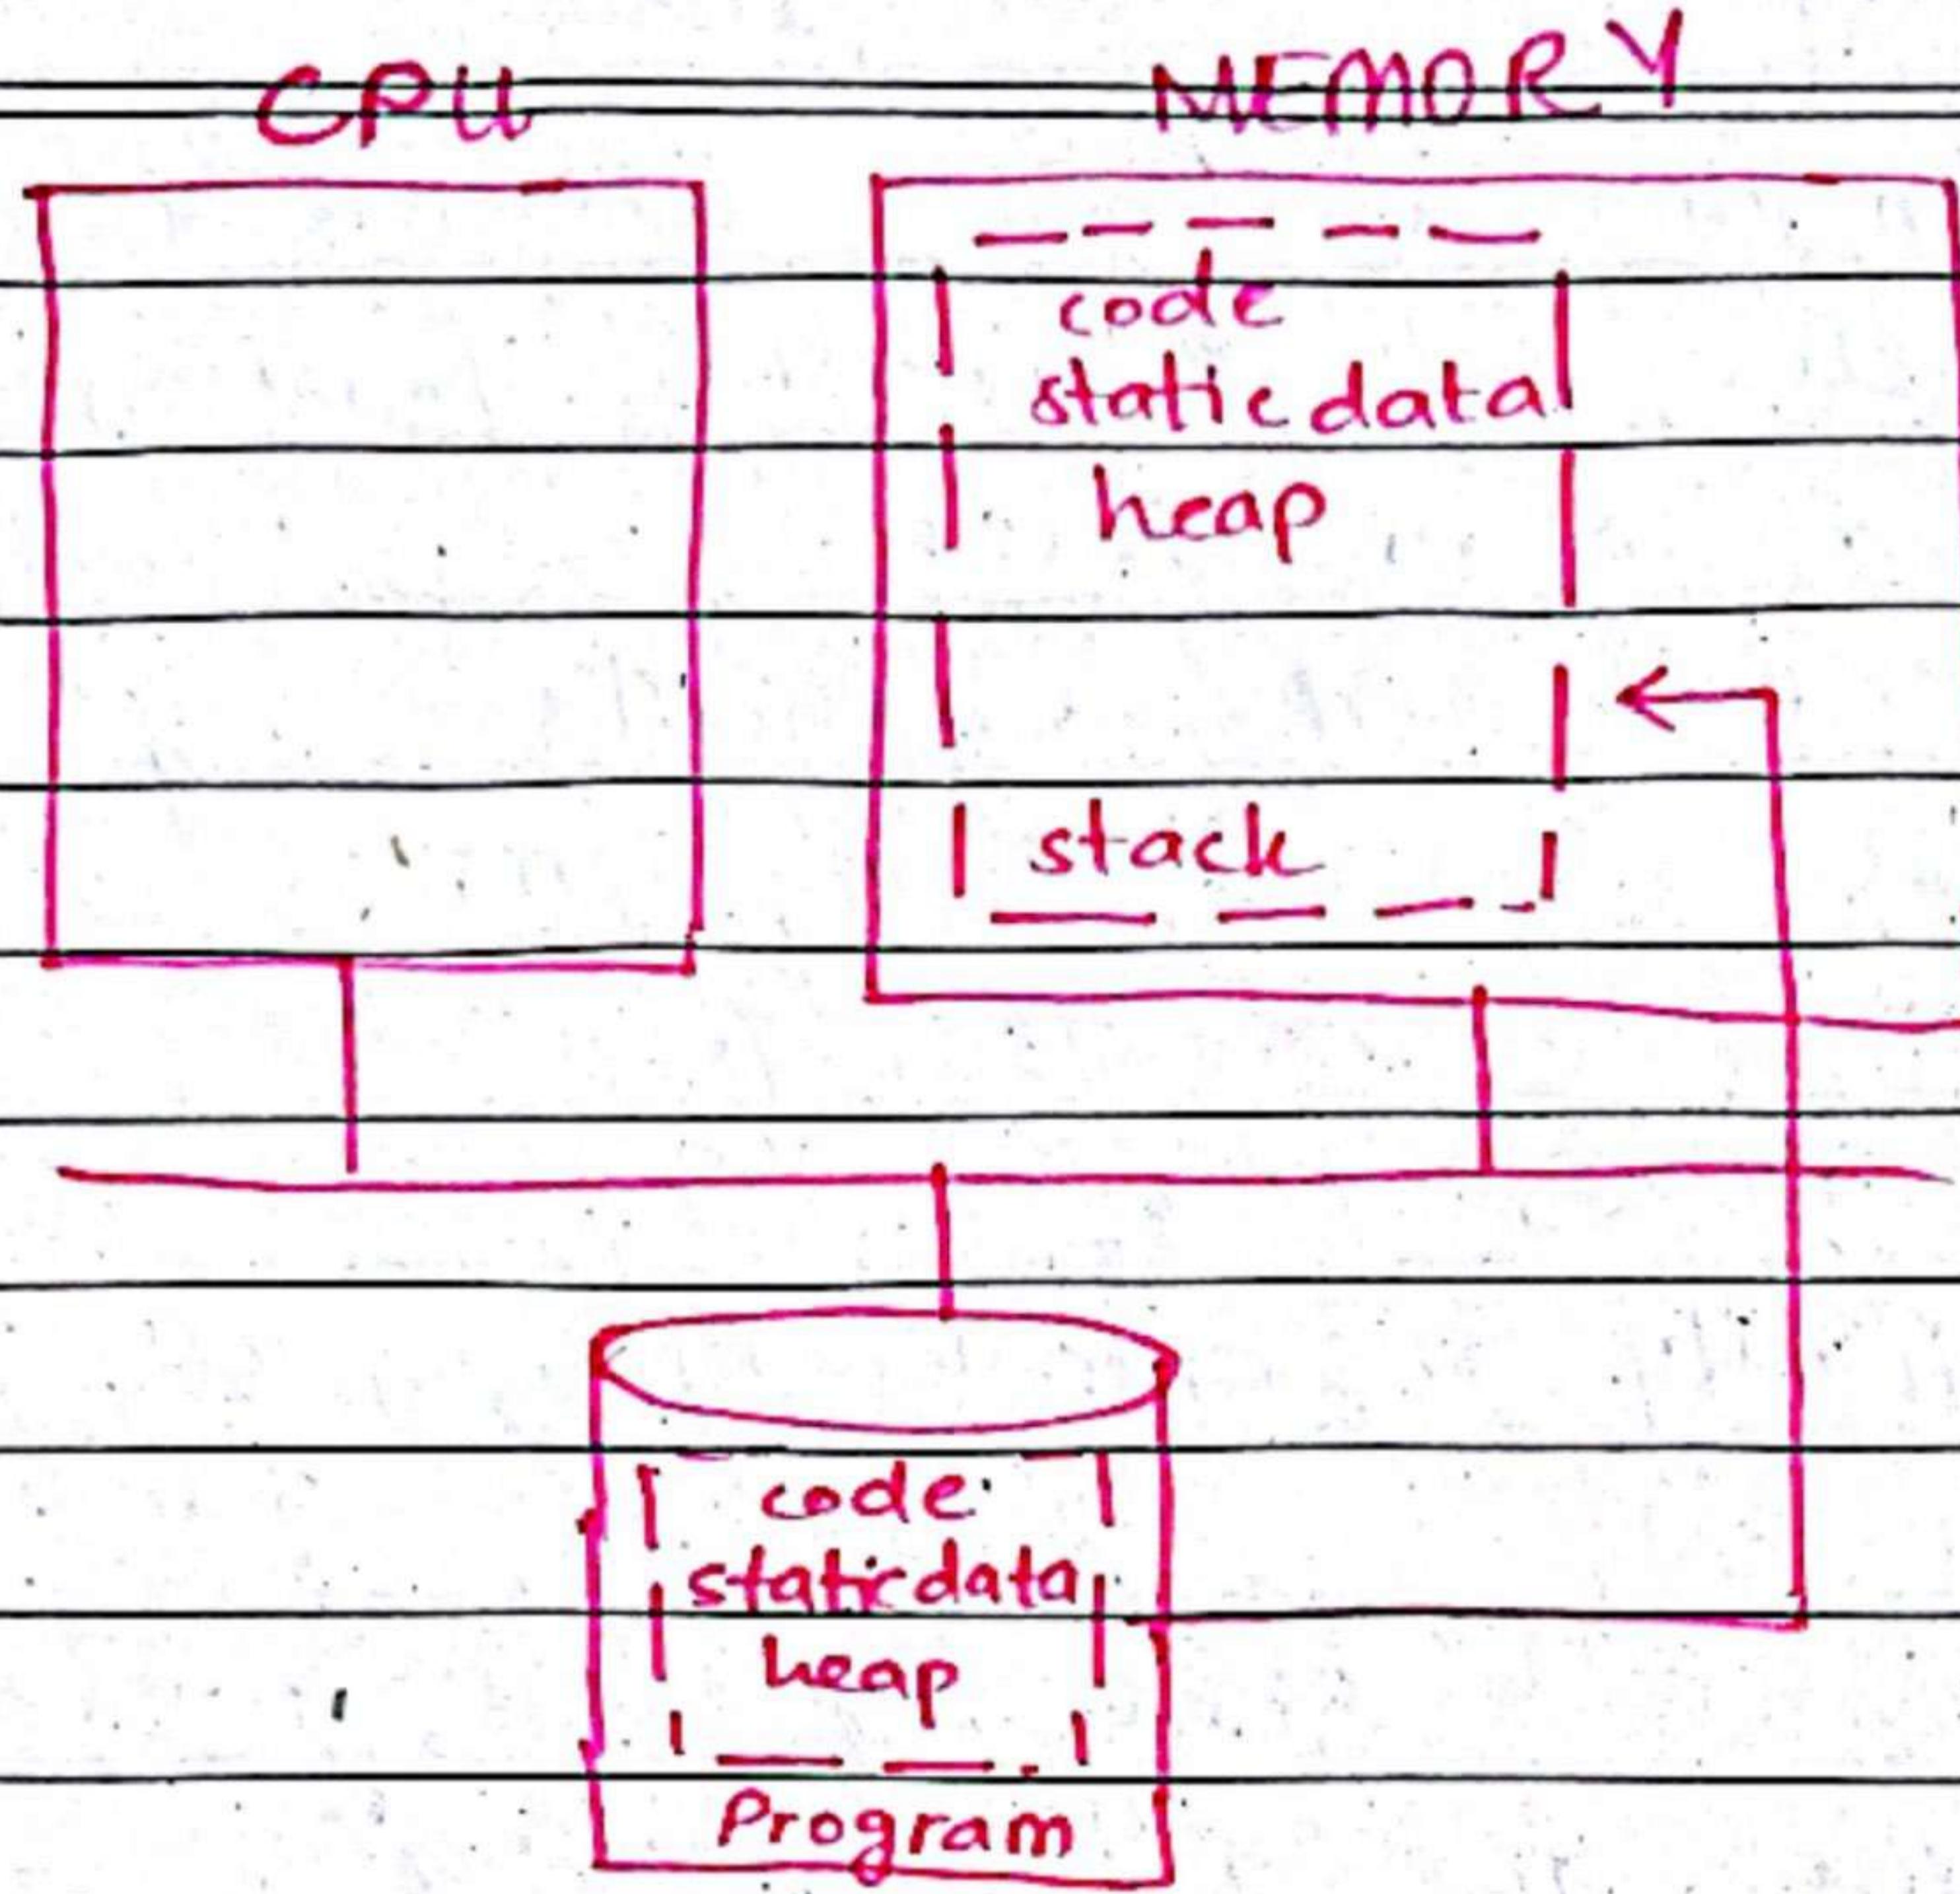
→ each process by default has three open file descriptors

→ standard input, output and error.

⑤ Start the program running at entry point, namely main()

→ OS transfers control of CPU to newly-created process.

Date: _____



loading:

Takes on-disk program
& reads it into address
space of process.

→ A process can be one of three states

→ Running

→ A process is running on a processor.

→ Ready

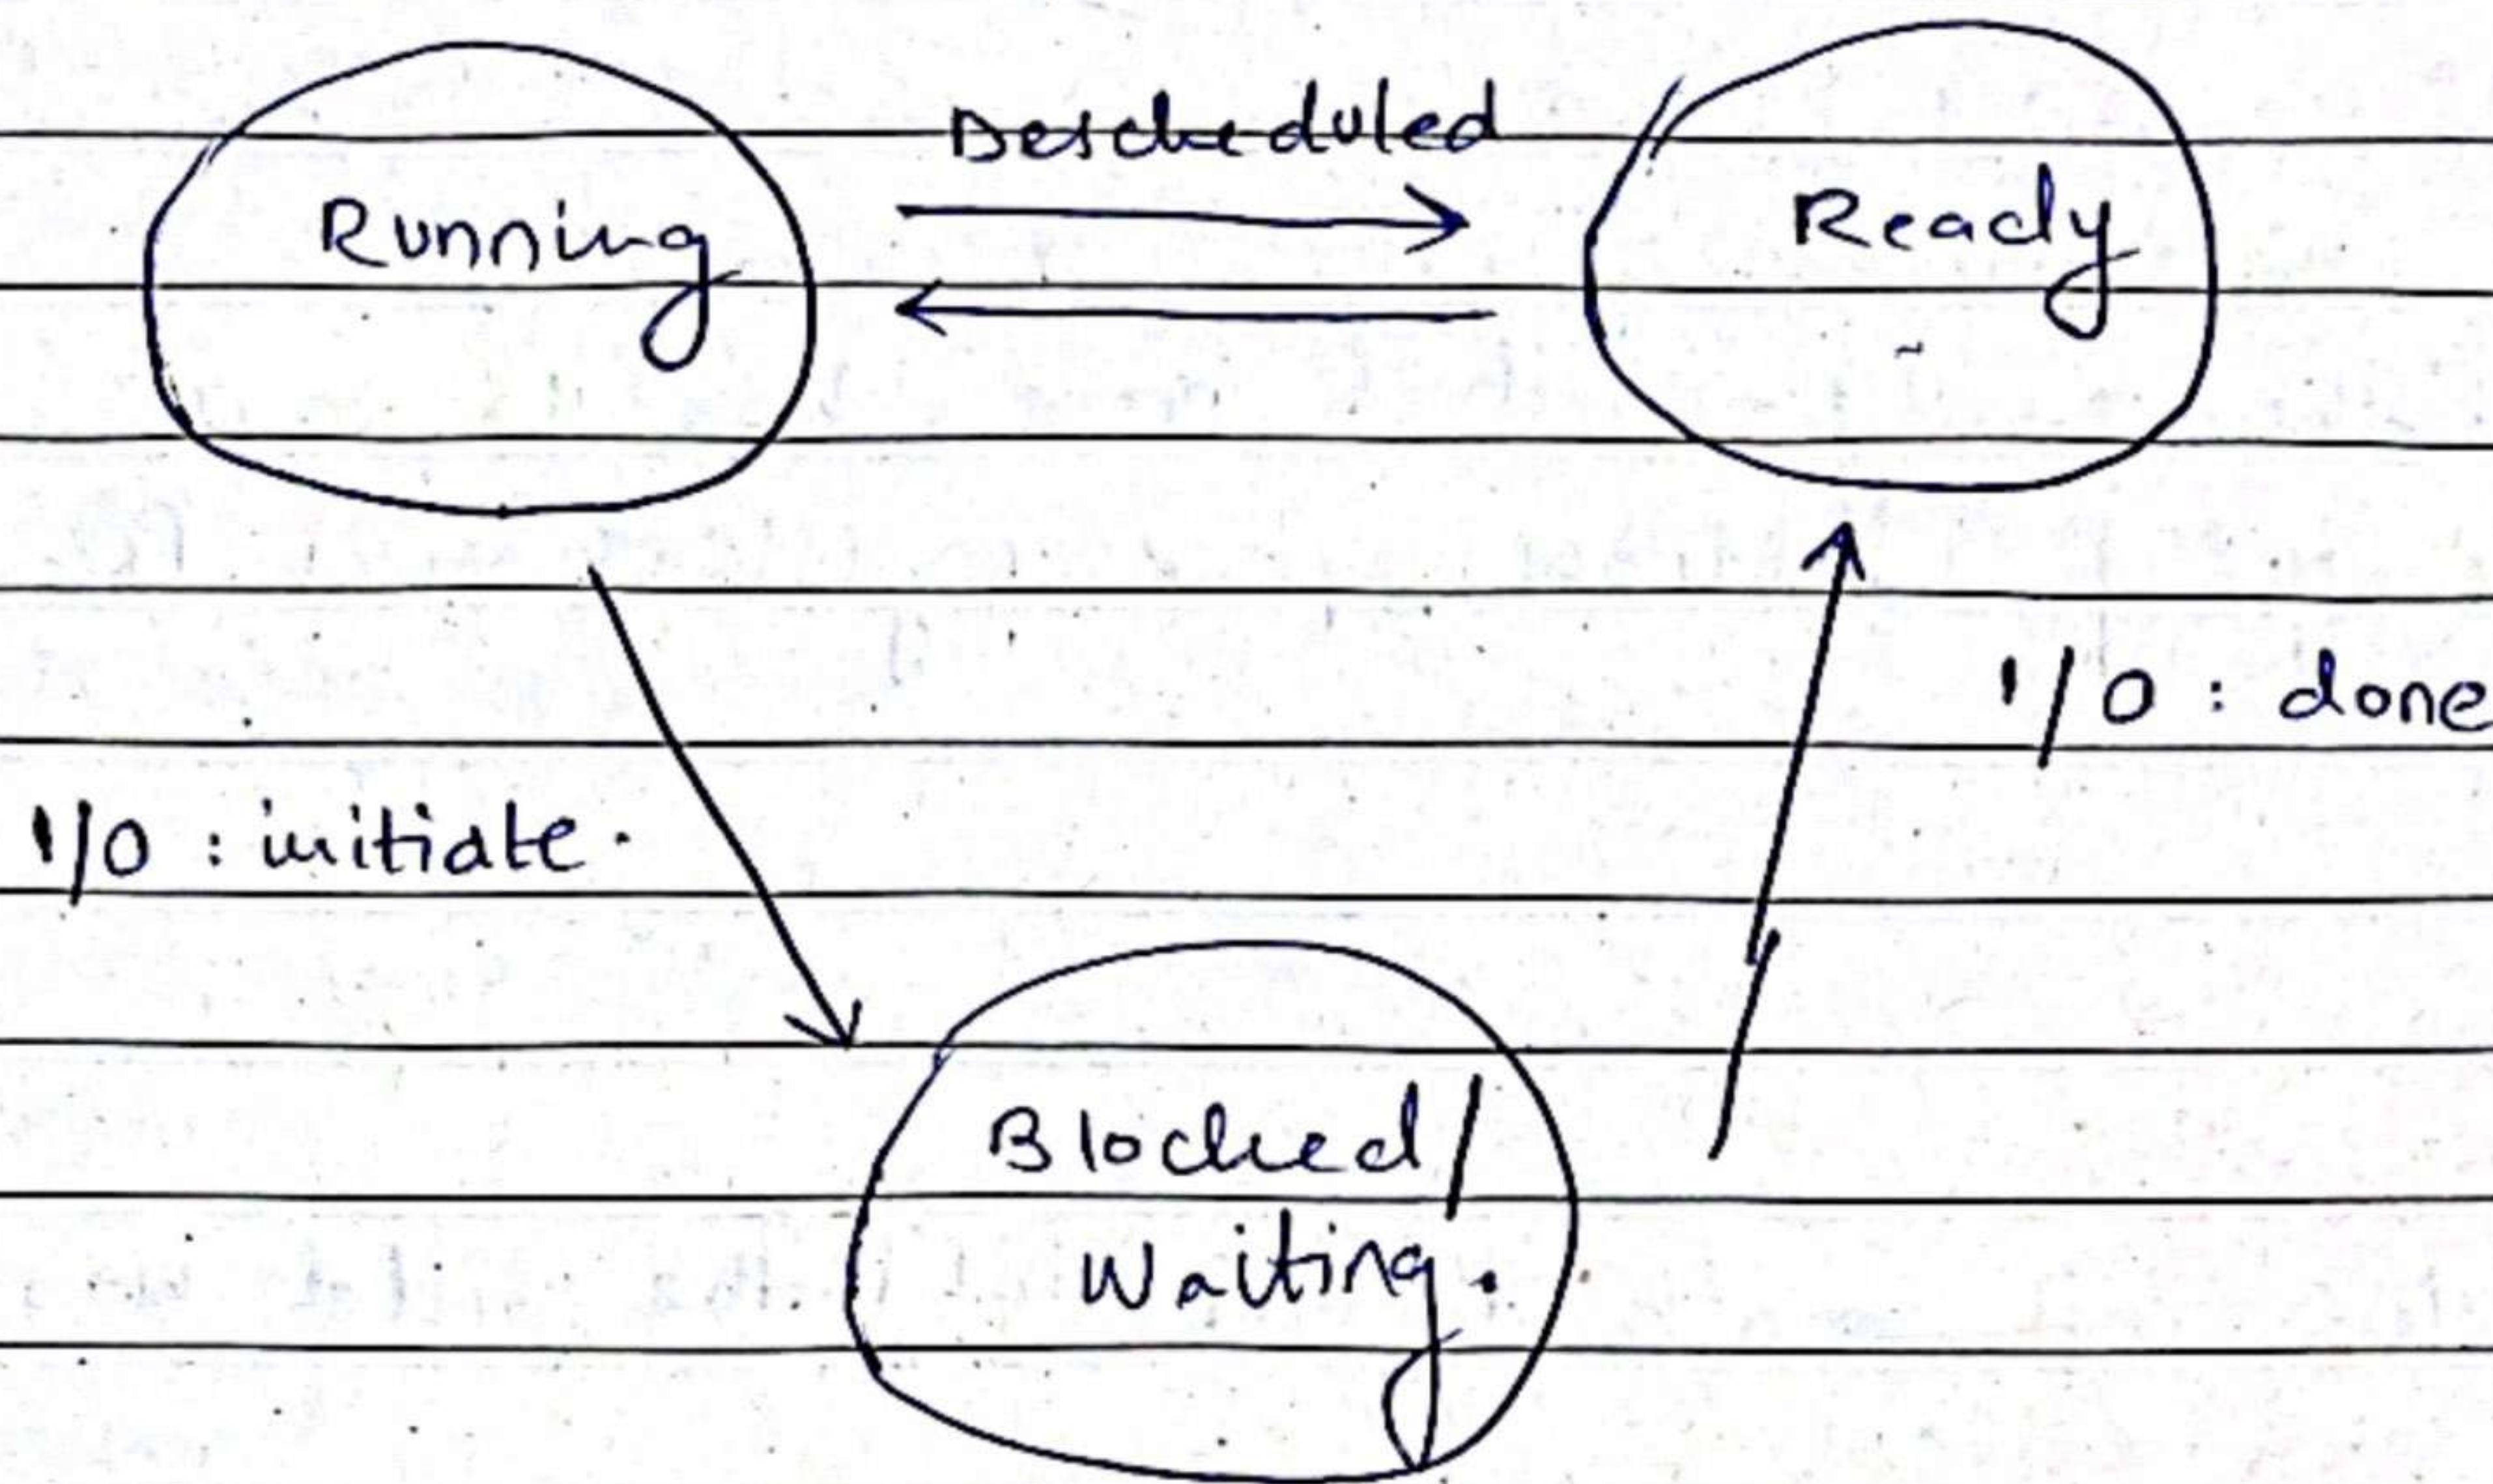
→ A process is ready to run but for some reason the OS has chosen not to run it at this given moment.

→ Blocked

→ A process has performed some kind of operation.

→ When a process initiates an I/O request to disk, it becomes blocked and thus some other process can use the processor.

a process in running state will go to ready if interrupt is raised. Date: _____



→ The OS has some key data structures that track various relevant pieces of information.

→ Process List

→ Ready processes.

→ Blocked process

→ Current running process

→ Register context.

→ PCB (Process Control Block)

→ A c-structure that contains information about each process.

* understanding of codes (refer to slide).

CHAPTER 5: Introduce : Process API → Date: _____

fork() system call

→ create new process

→ the newly-created process has its own copy of address space, registers and PC.

wait() system call

→ this system call won't return until the child has run and exited.

exec() system call

→ run a program that is different from calling program.

CHAPTER 6: Mechanism : Limited Direct Execution.

How to efficiently virtualize the CPU with control?

→ The OS needs to share the physical CPU by time sharing

→ Issue

- Performance: How can we implement virtualization without adding excessive overhead to the system?

- Control: How can we run processes efficiently while retaining control over CPU?

Date: _____

DIRECT EXECUTION

→ Just run the program directly on the CPU.

OS

Programs

- ① Create entry for process list
- ② Allocate memory for program
- ③ Load program into memory
- ④ Set up stack with argc/argv
- ⑤ Clear registers
- ⑥ Execute call `main()`
- ⑦ Run `main()`
- ⑧ Execute return from `main()`
- ⑨ Free memory of process
- ⑩ Remove from process list

Without limits on running programs, the OS wouldn't be in control of anything and thus would be "just library"

An interrupt within OS occurs 3 times:

- ① Program misbehaves falls into segmentation fault.
- ② During input/output.

- ③ system calls.

→ BIOS loads OS and also loads
Interrupt Disrupt Table (IDT)

→ Interrupt Disrupt Table (IDT) → IDT knows where ~~Data~~ system call is.

PROBLEM 1 Restricted Operation.

→ What if a process wishes to perform some kind of restricted operation such as ...

→ issuing an I/O request to disk

→ gaining access to more system resources such as CPU or a memory.

SOLUTION

Using protected control transfer.

→ **User Mode:** Applications do not have full access to hardware resources.

→ **Kernel Mode:** The OS has full access to the full resources of machine.

SYSTEM CALL

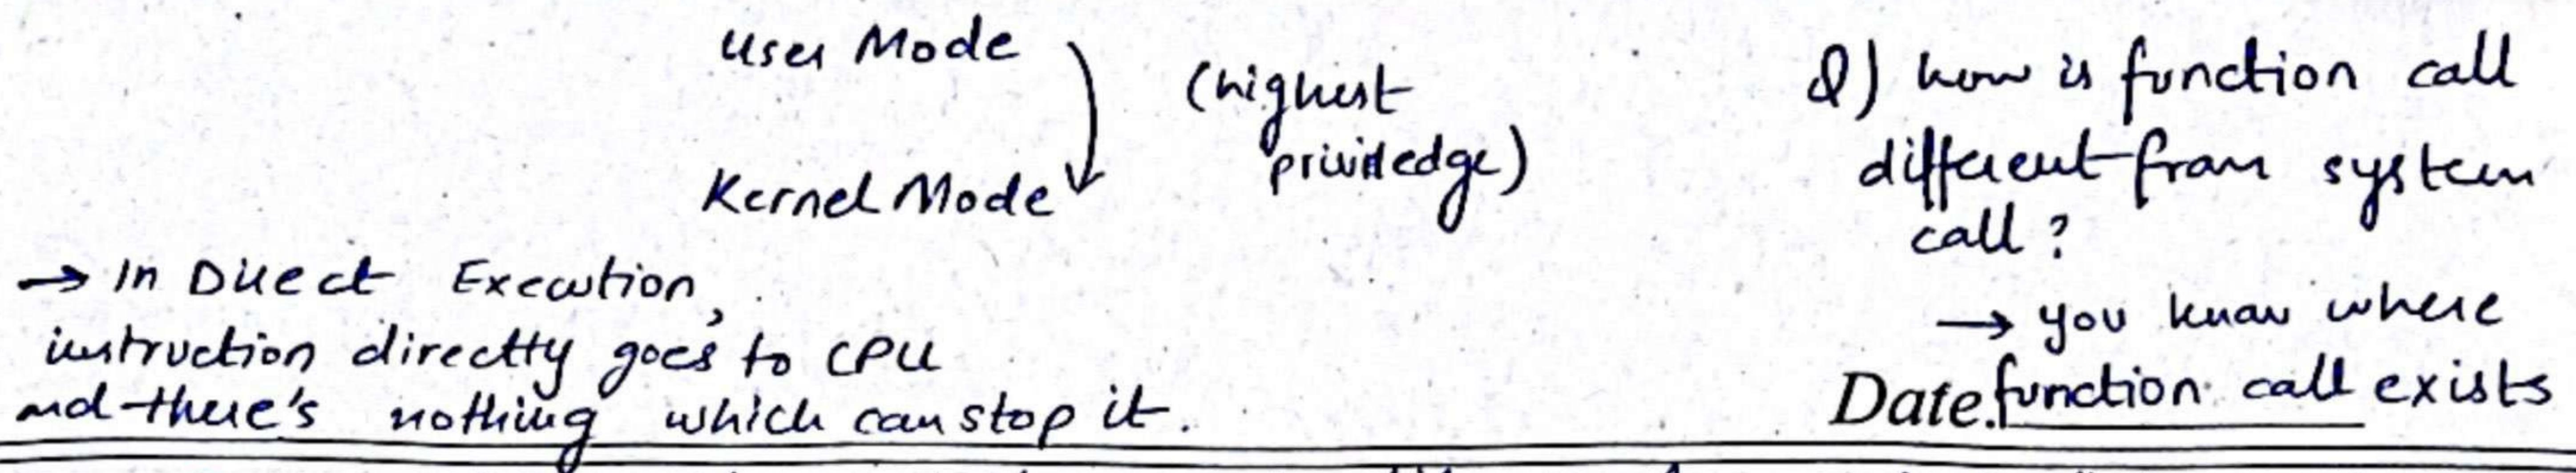
→ Allow the kernel to carefully expose certain key pieces of functionality to user program such as ...

↳ accessing the filesystem

↳ creating and destroying processes

↳ communicating with other processes

↳ allocating more memory.



Q) how is function call different from system call?
→ you know where Date.function call exists

TRAP → os is right now just a massive library of functions.

but don't know where system call exists.

→ jump into the kernel.

→ raise the privilege level to kernel mode

RETURN-FROM-TRAP

→ return into calling user program.

→ reduce the privilege level back to user mode.

PROBLEM 2 switching Between Processes

→ How can the OS regain control of the CPU so that it can switch b/w processes?

→ a cooperative approach : Wait for system calls

→ a non-cooperative approach : The OS takes control.

COOPERATIVE APPROACH (when control is given to OS)

→ Processes periodically give up the CPU by making system calls such as yield

→ The OS decides to run some other task.

→ Application also transfer control to the OS when they do something illegal.

→ Divide by zero

→ try access memory that it shouldn't be able to access.

A process gets stuck in an infinite loop

→ Reboot the machine.

Date: _____

NON-COOPERATIVE APPROACH

- A timer interrupt
 - during boot sequence, the OS starts the timer.
 - The timer raises an interrupt every so many milliseconds.
- When the interrupt is raised:
 - The currently running process is halted
 - Save enough of the state of the program.
 - A pre-configured interrupt handler in OS runs.

A timer interrupt gives OS the ability to run again on a CPU.

SCHEDULER (Saving & Restoring Context)

- Scheduler makes decision
 - whether to continue running the current process or switch to different one.
- if the decision is made to switch, the OS executes context switch

→ save pre register
of old process and load
new

Date: _____

CONTEXT SWITCH

- A low-level piece of assembly code.
- Save a few register values for the current process onto its kernel stack
 - general purpose registers
 - PC
- Kernel Stack Pointer
- Restore a few for the soon-to-be-executing process from its kernel stack
- Switch to the kernel stack for soon-to-be-executing process.

CHAPTER 7: Scheduling.

Purpose of Scheduling.

Manges execution of processes or threads on a CPU to optimize certain performance metrics.

Goals

- ① Minimize CPU utilization
- ② Minimise response time
- ③ Minimize turnaround time.

Workload assumptions

same

- each job runs for [^]amount of time.
- all jobs arrive at same time.
- all jobs only use CPU (ie they perform no I/O)
- the run-time of each job is known.

* time from when the job arrives to first time it is scheduled.

$$T_{\text{response}} = T_{\text{first run}} - T_{\text{arrival}}$$

* Decreasing Response Time increases turnaround time.

Date: _____

- Turnaround Time: Total time taken from submission to completion of process.
- Response Time: Time from submission to first response produced.
- Throughput: Number of processes completed per time unit

$$T_{\text{turnaround}} = T_{\text{completion}} - T_{\text{arrival}}$$

* the time at which the job completes minus the time at which the job arrived in the system.

→ Another metric is fairness

→ performance and fairness are often at odds in scheduling.

SCHEDULING ALGORITHMS

① First Come, First Served (FCFS)

→ very simple & easy to implement.

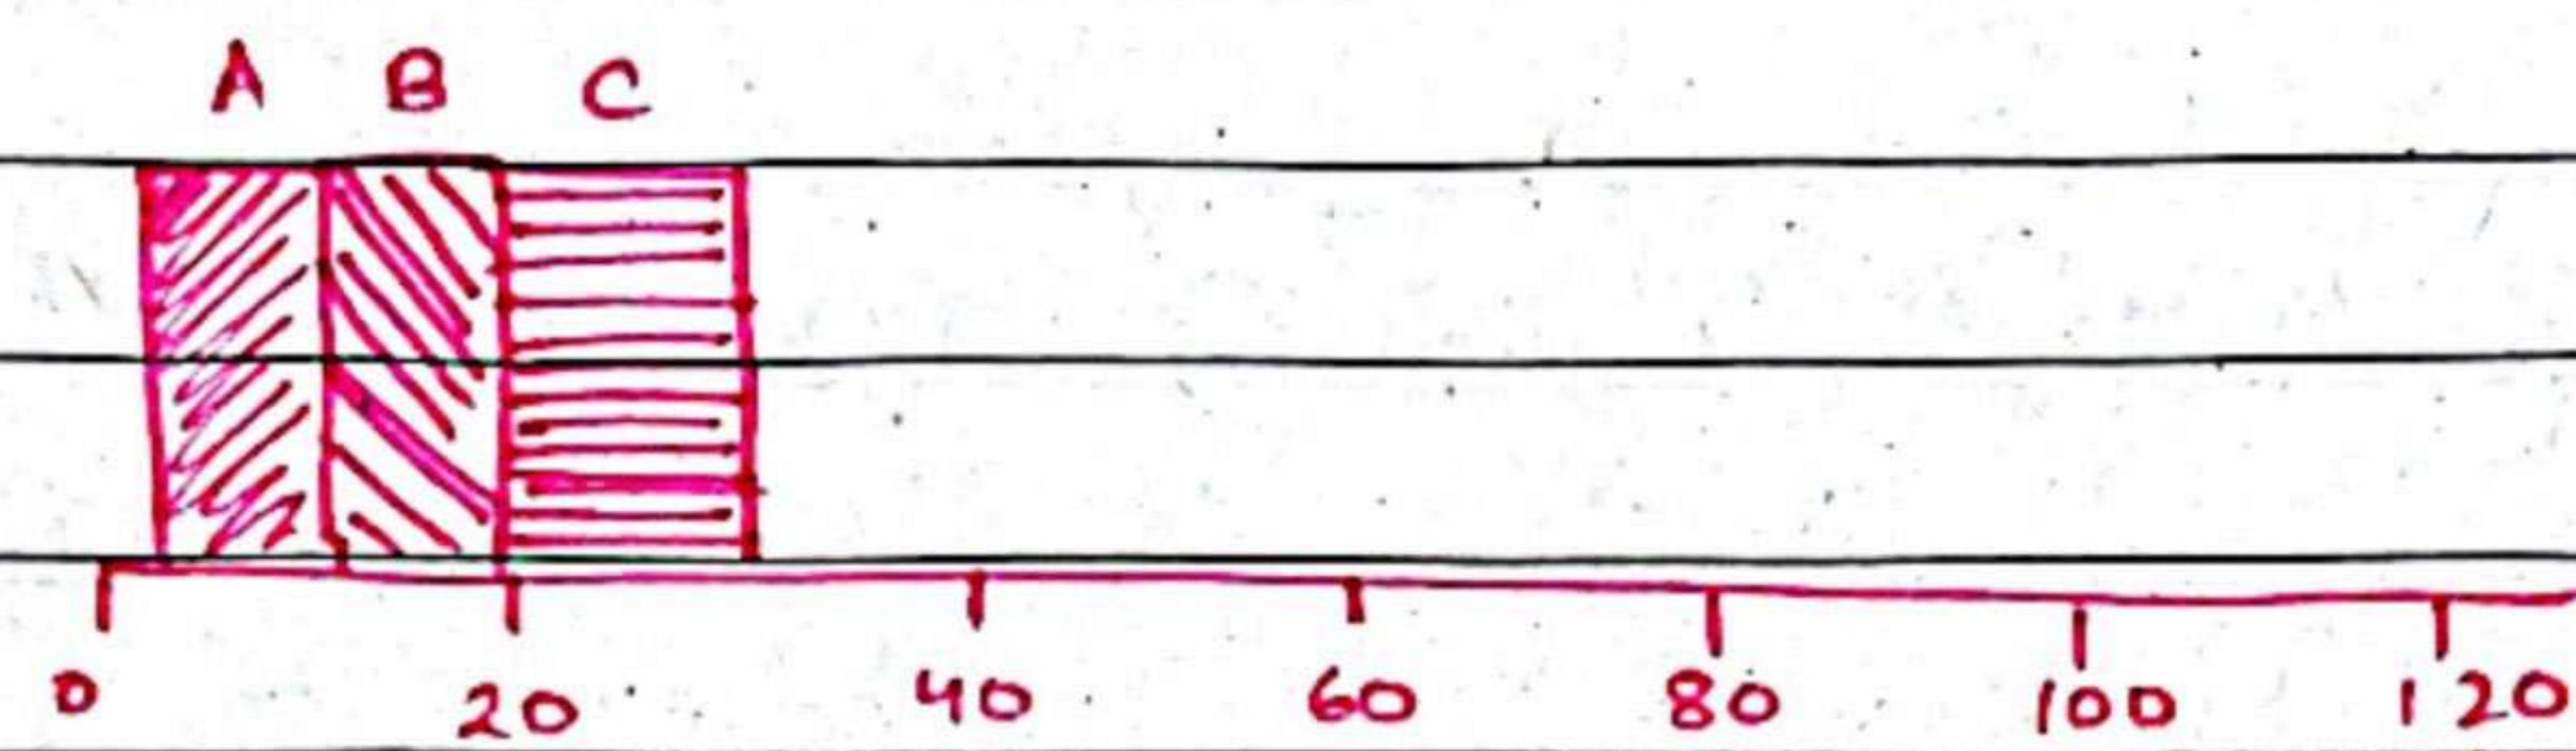
→ processes are scheduled in the order they arrive.

Example

→ A arrived just before B which arrived just before C.

→ each job runs for 10 seconds.

Date: _____



$$\text{Average turnaround time} = \frac{10 + 20 + 30}{3} = 20 \text{ sec.}$$

→ Let's relax assumption 1: Each job no longer runs for the same amount of time.

Example

- A just arrived before B which arrived just before C.
- A runs for 100 seconds, B and C runs for 10 each.

$$\text{Average turnaround time} = \frac{100 + 110 + 120}{3} = 110 \text{ sec.}$$

② Shortest Job First (SJF)

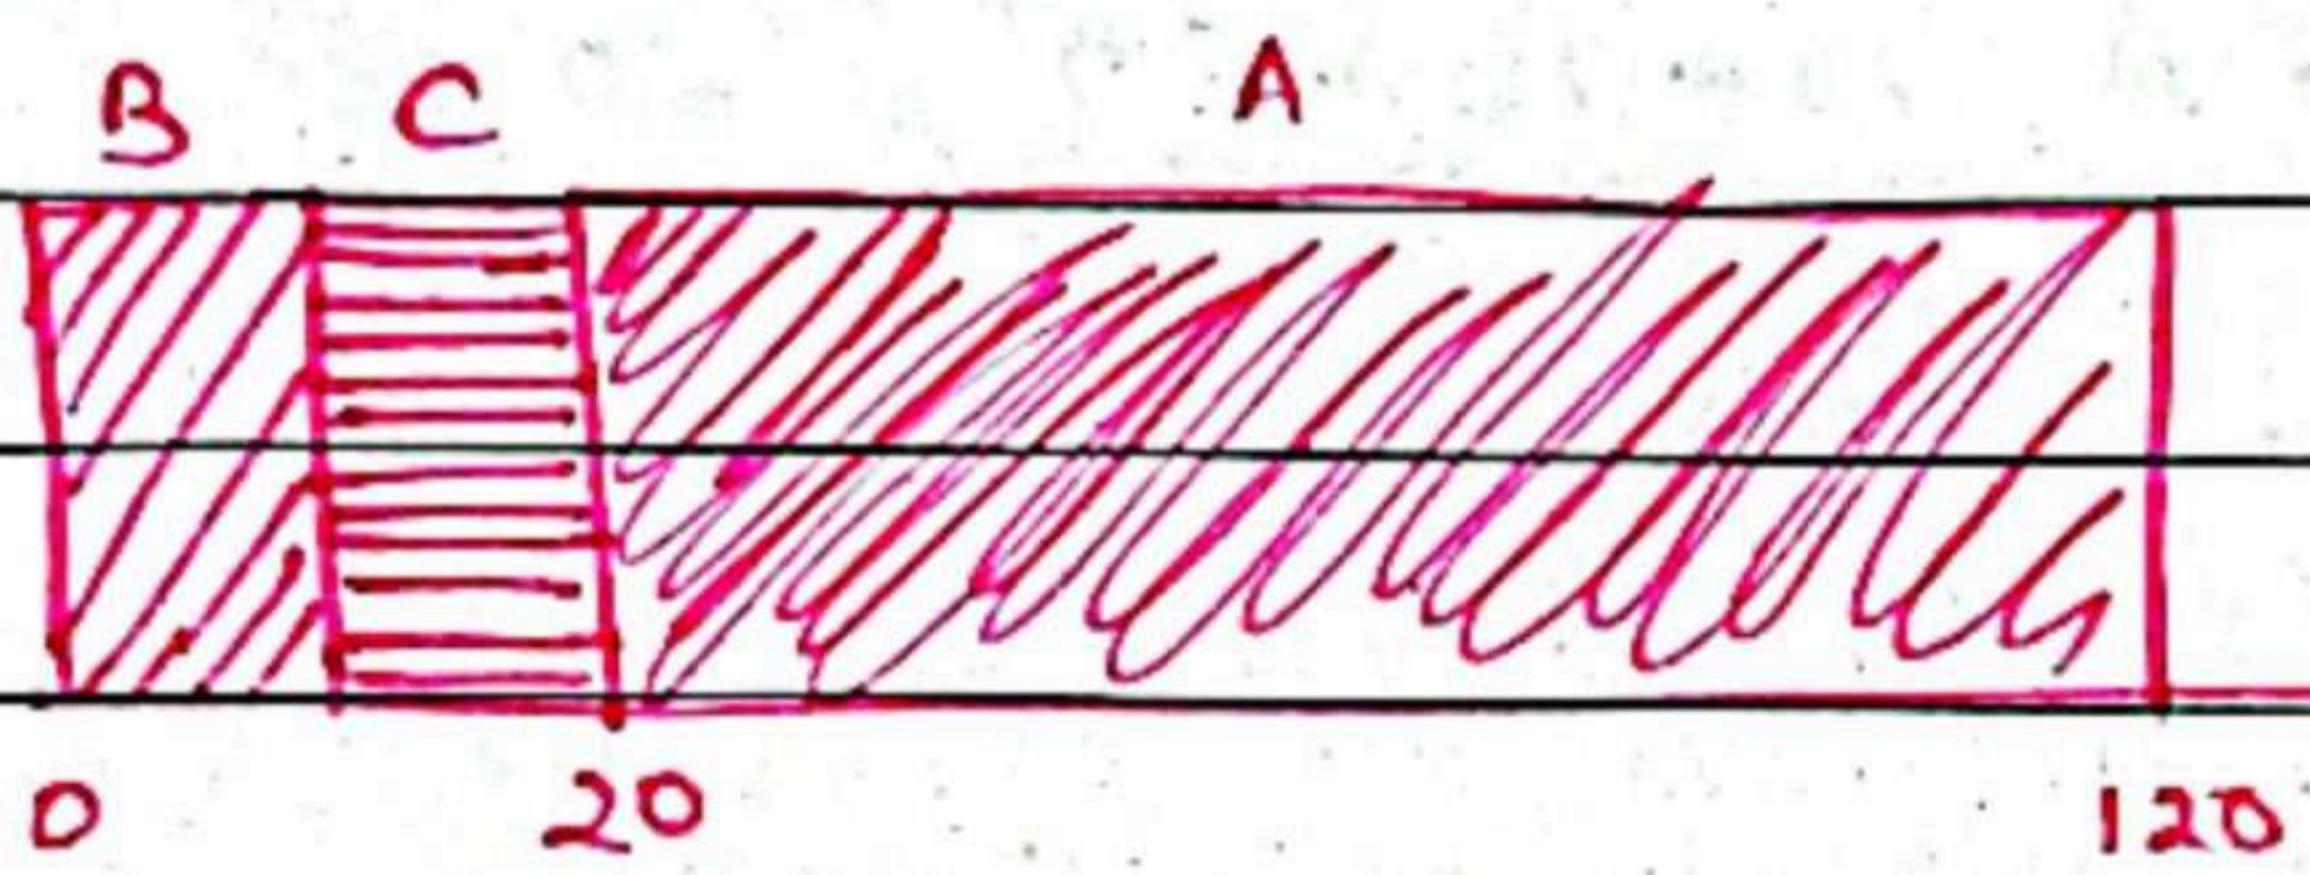
- executes the shortest upcoming job next.
- run the shortest job first, then next shortest and so on.
- NON-PREEMPTIVE SCHEDULER.

Date: _____

Example

→ A arrived just before B which arrived just before C

→ A runs for 100 seconds, B and C run for 10 each.



$$\text{Turnaround time} = \frac{10 + 20 + 120}{3} = 50 \text{ sec.}$$

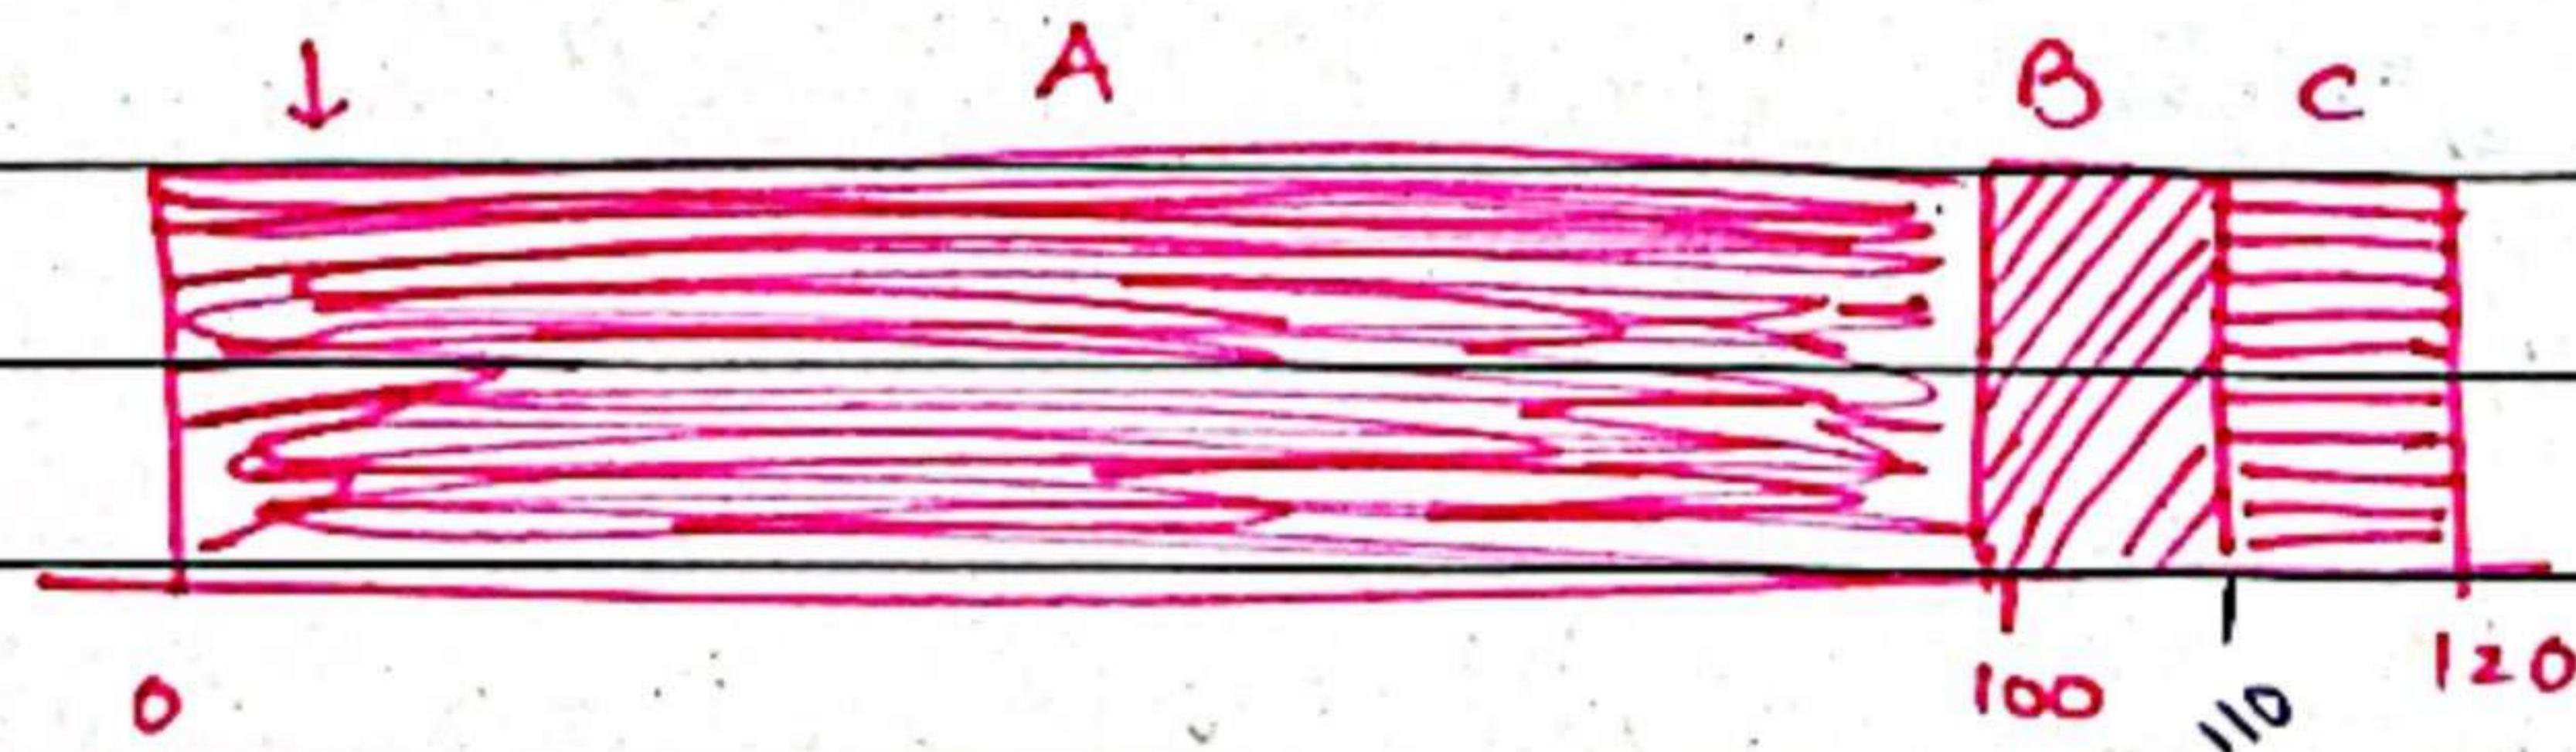
→ Let's relax assumption 2: Jobs can arrive at any time

Example:

→ A arrives at $t=0$ and needs to run for 100 seconds.

→ B and C arrive at $t=10$ and each need to run for 10 seconds.

[B, C arrive]



* STCF optimizes turnaround time, RR optimizes response time.

Date:

$$\text{Average turnaround time} = \frac{100 + (110-10) + (120-10)}{3} = 103.33 \text{ sec}$$

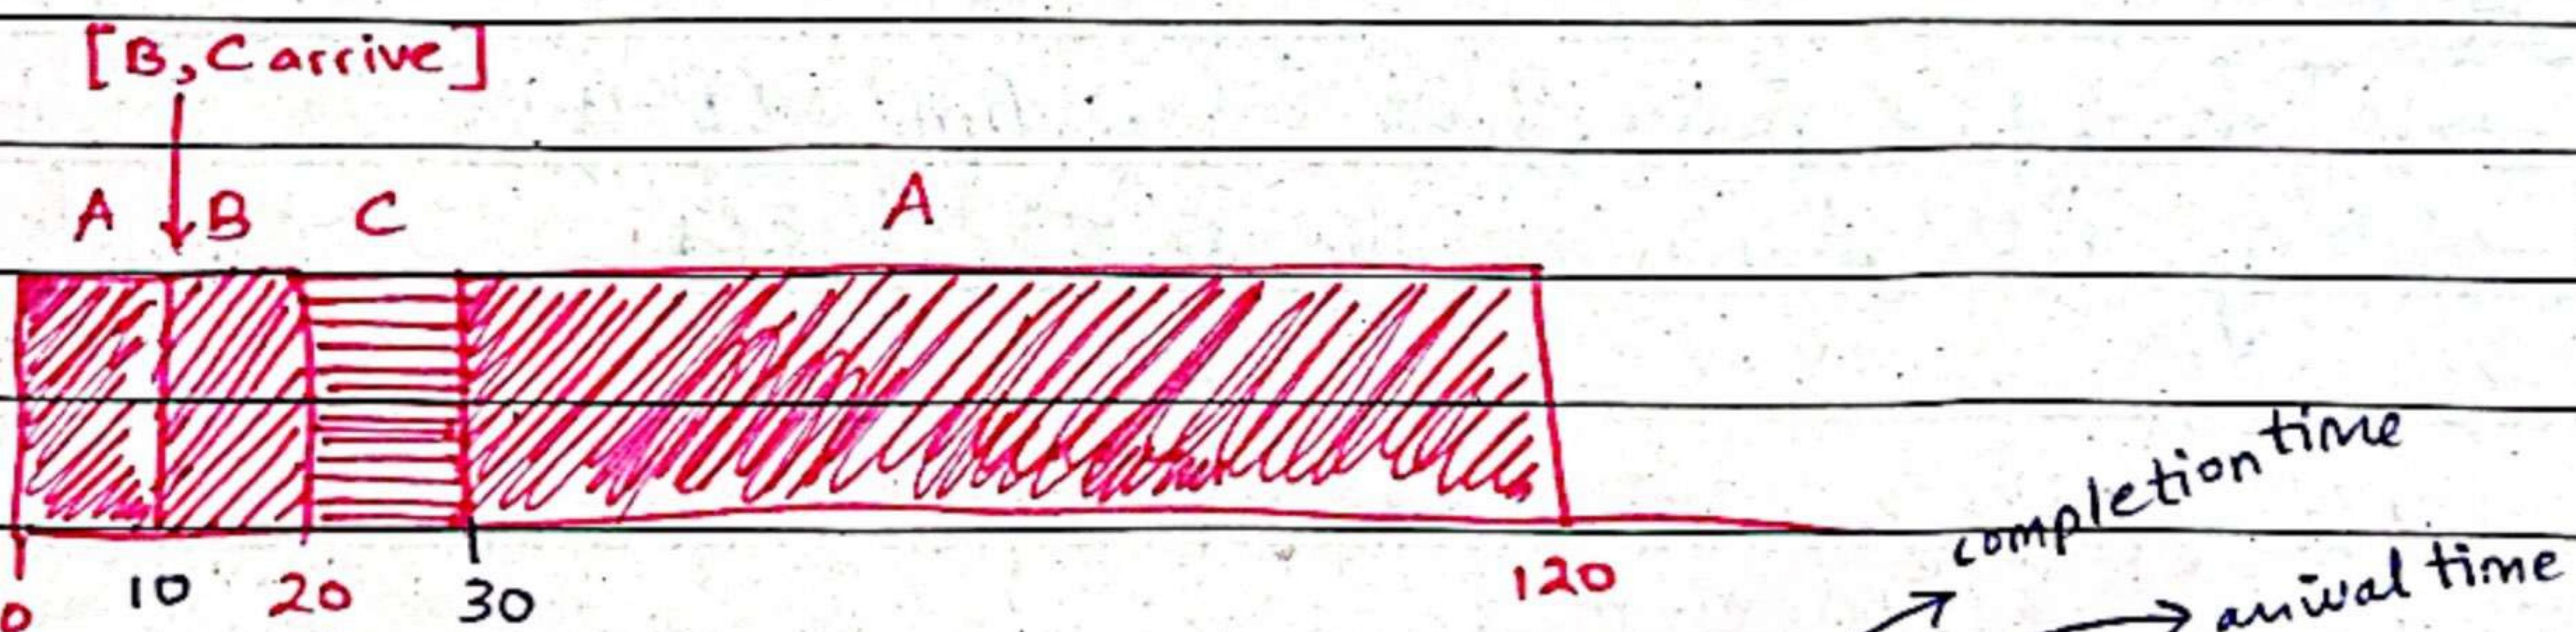
③ Shortest Time-to-completion First (STCF)

→ Preemptive version of SJF, where current running process can be interrupted.

Example

→ A arrives at $t=0$ and needs to run for 100 seconds.

→ B and C arrive at $t=10$ and each need to run for 10 seconds.



$$\text{Average turnaround time} = \frac{(120-0) + (20-10) + (30-10)}{3} = 50 \text{ sec.}$$

④ Round Robin (RR) Scheduling

→ each process receives a fixed time slot, rotating through queue.

Date: _____

→ Time Slicing

→ run job for a time slice and then switch to next job in the run queue until the jobs are finished.

→ time slice is sometimes called scheduling quantum

→ it repeatedly does so until the jobs are finished

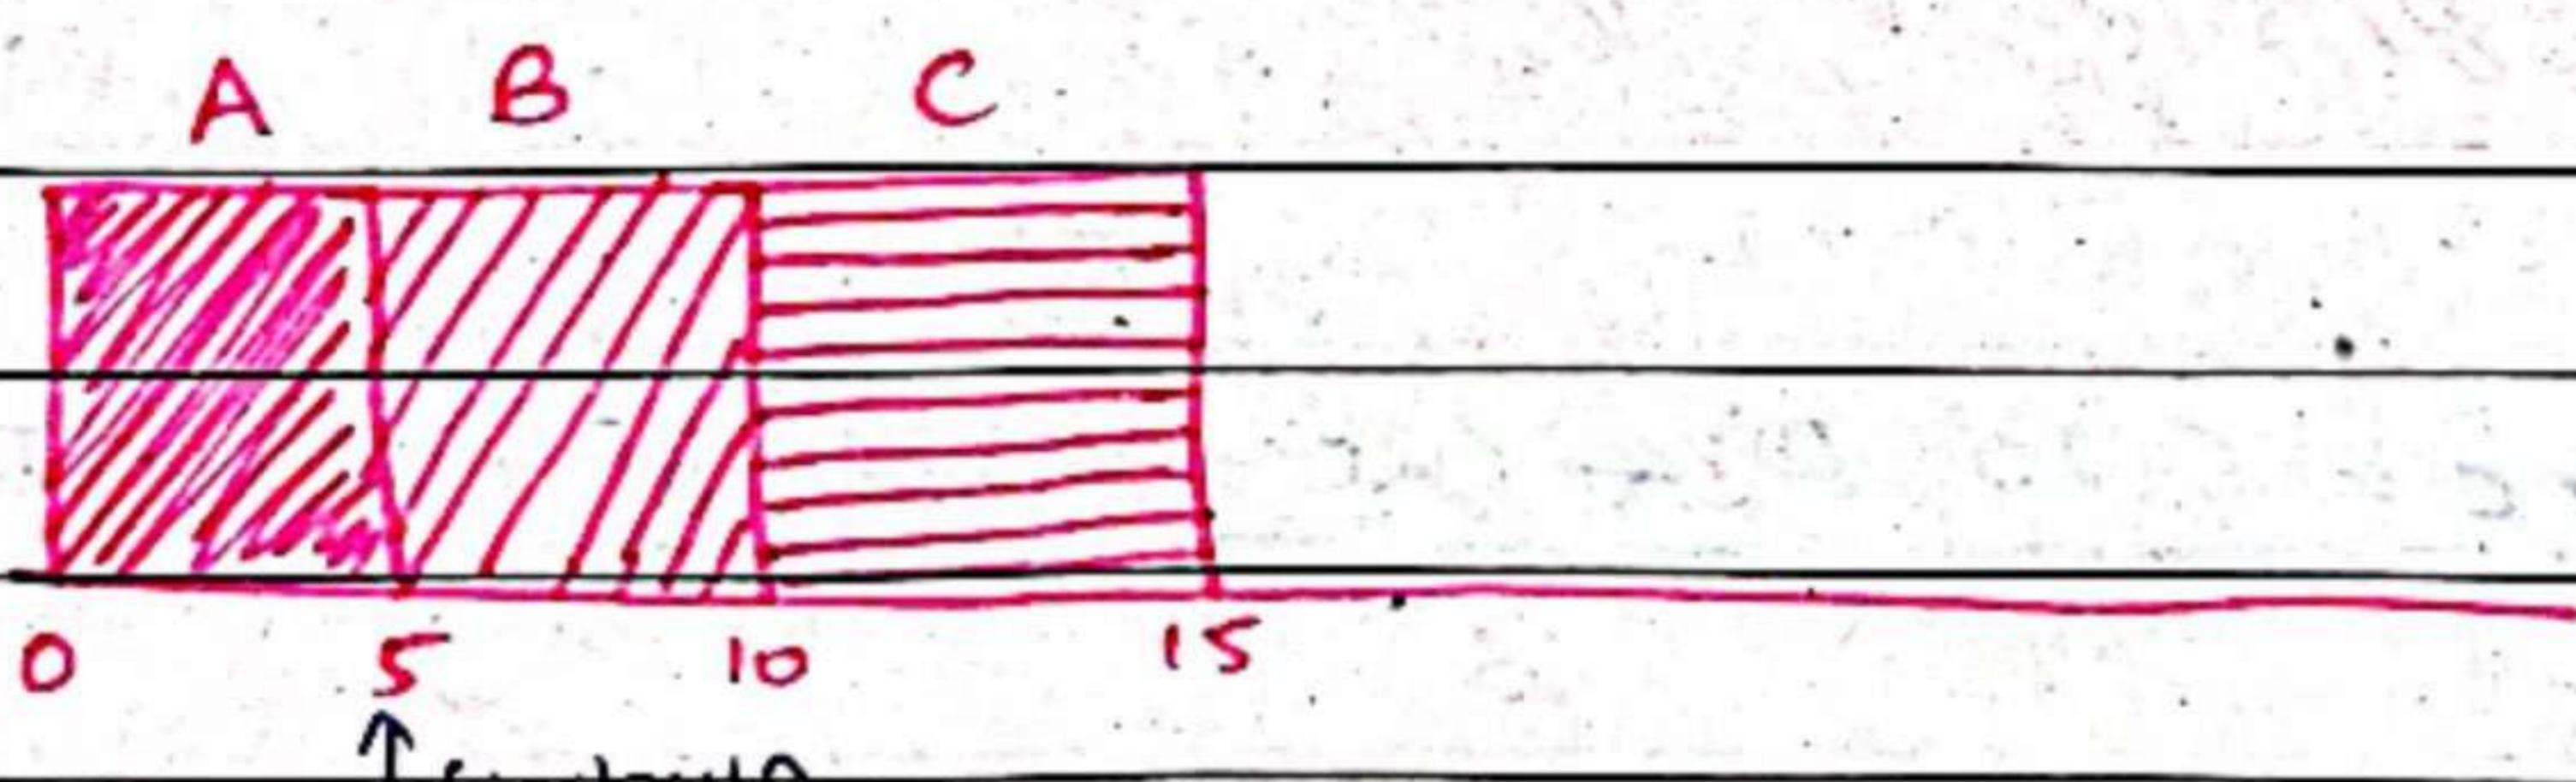
→ The length of time slice must be multiple of time-interrupt period.

RR is fair but performs poorly on metrics such as turnaround time.

Example

→ A, B and C [arrive] at same time. ($t=0$)

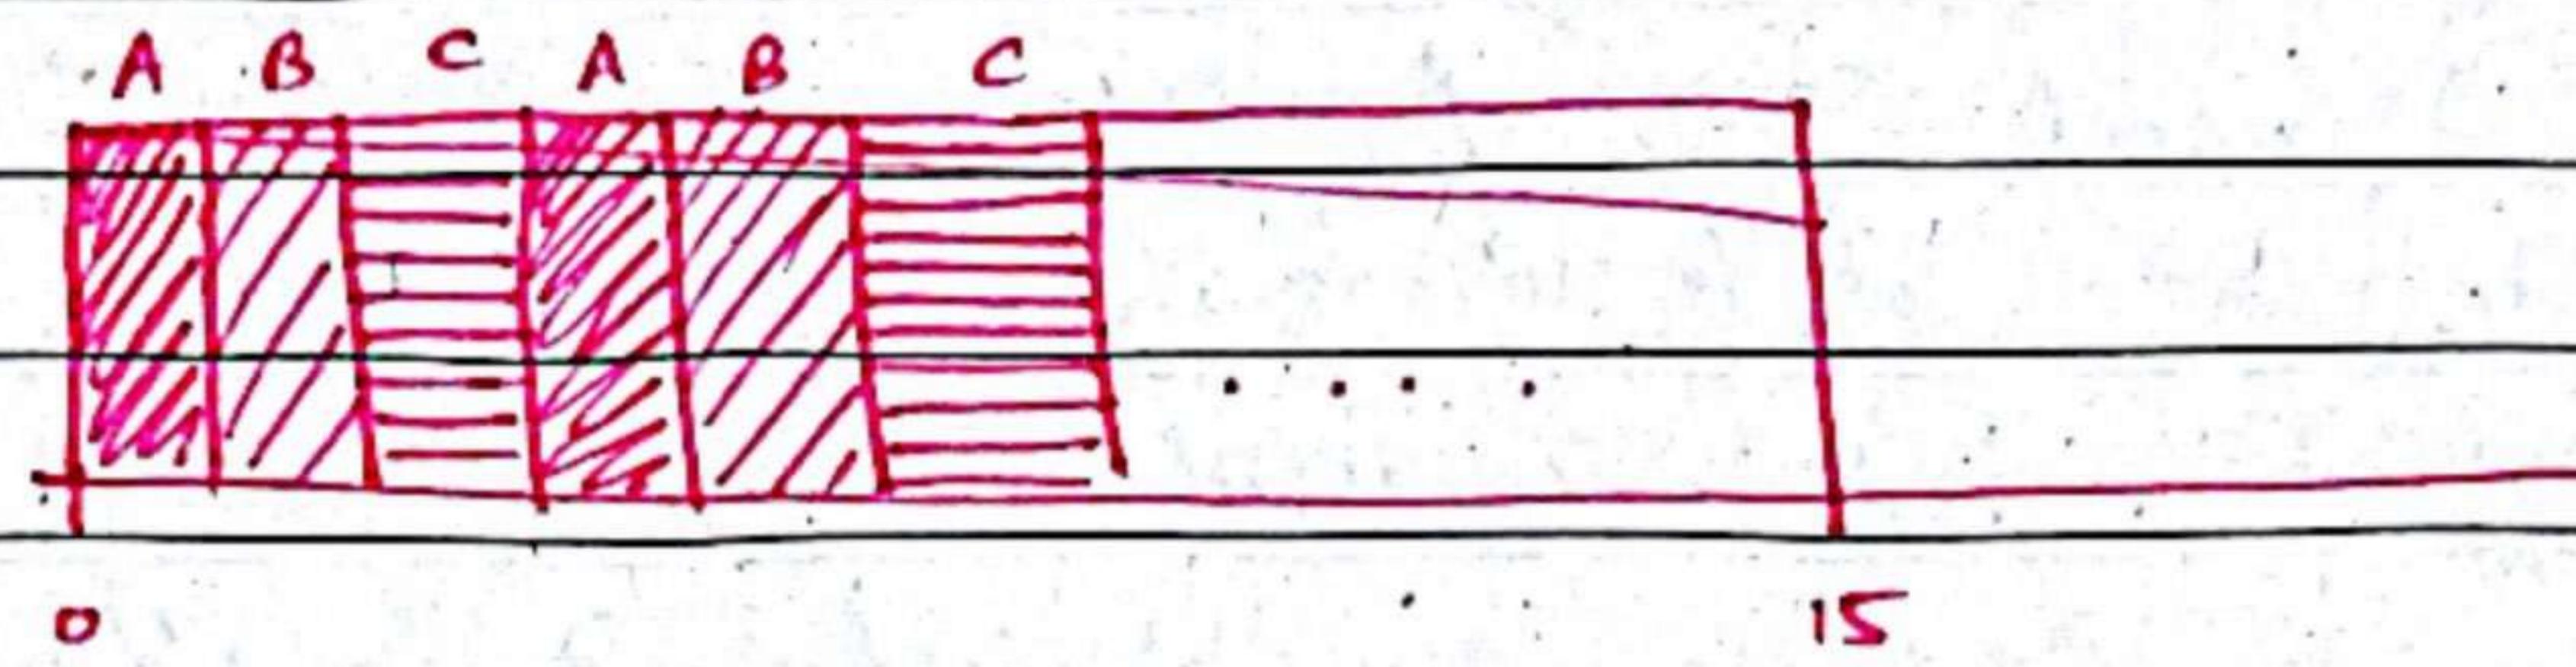
→ they each wish to run for 5 seconds.



$$T_{\text{average response}} = \frac{0 + 5 + 10}{3} = 5 \text{ sec.}$$

$$= T_{\text{fustrun}} - T_{\text{arrival}}$$

Date: _____



$$T_{\text{average response}} = \frac{0+1+2}{3} = 1 \text{ sec.}$$

(RR with time-slice of 1 sec - Good for Response Time)

The length of time slice is critical.

- The shorter time slice.
 - better response time,
 - the cost of context switching will dominate overall performance
- The larger time slice.
 - amortize the cost of switching
 - worse response time

Deciding on length of time slice presents a trade-off to a system designer.

Date: _____

→ let's relax assumption 3: All programs perform I/O

Example

- A and B need 50ms of CPU time each.
- A runs for 10ms and then issues an I/O request
 - I/O ~~task~~ each take 10ms.
- B simply uses CPU for 50ms & performs no I/O.
- The scheduler runs A first, then B after.

CHAPTER 8: Scheduling: Multilevel Feedback Queue.

Multi-level Feedback Queue (MLFQ)

- A complex scheduler designed to optimize both turnaround time & response time without prior knowledge of job lengths.
- In simple words, a scheduler that learns from the past to predict future
- Objective:
 - ① Optimize turnaround time → run shorter jobs first
 - ② Minimize response time ~~&~~ without a prior knowledge of job length.

Date: _____

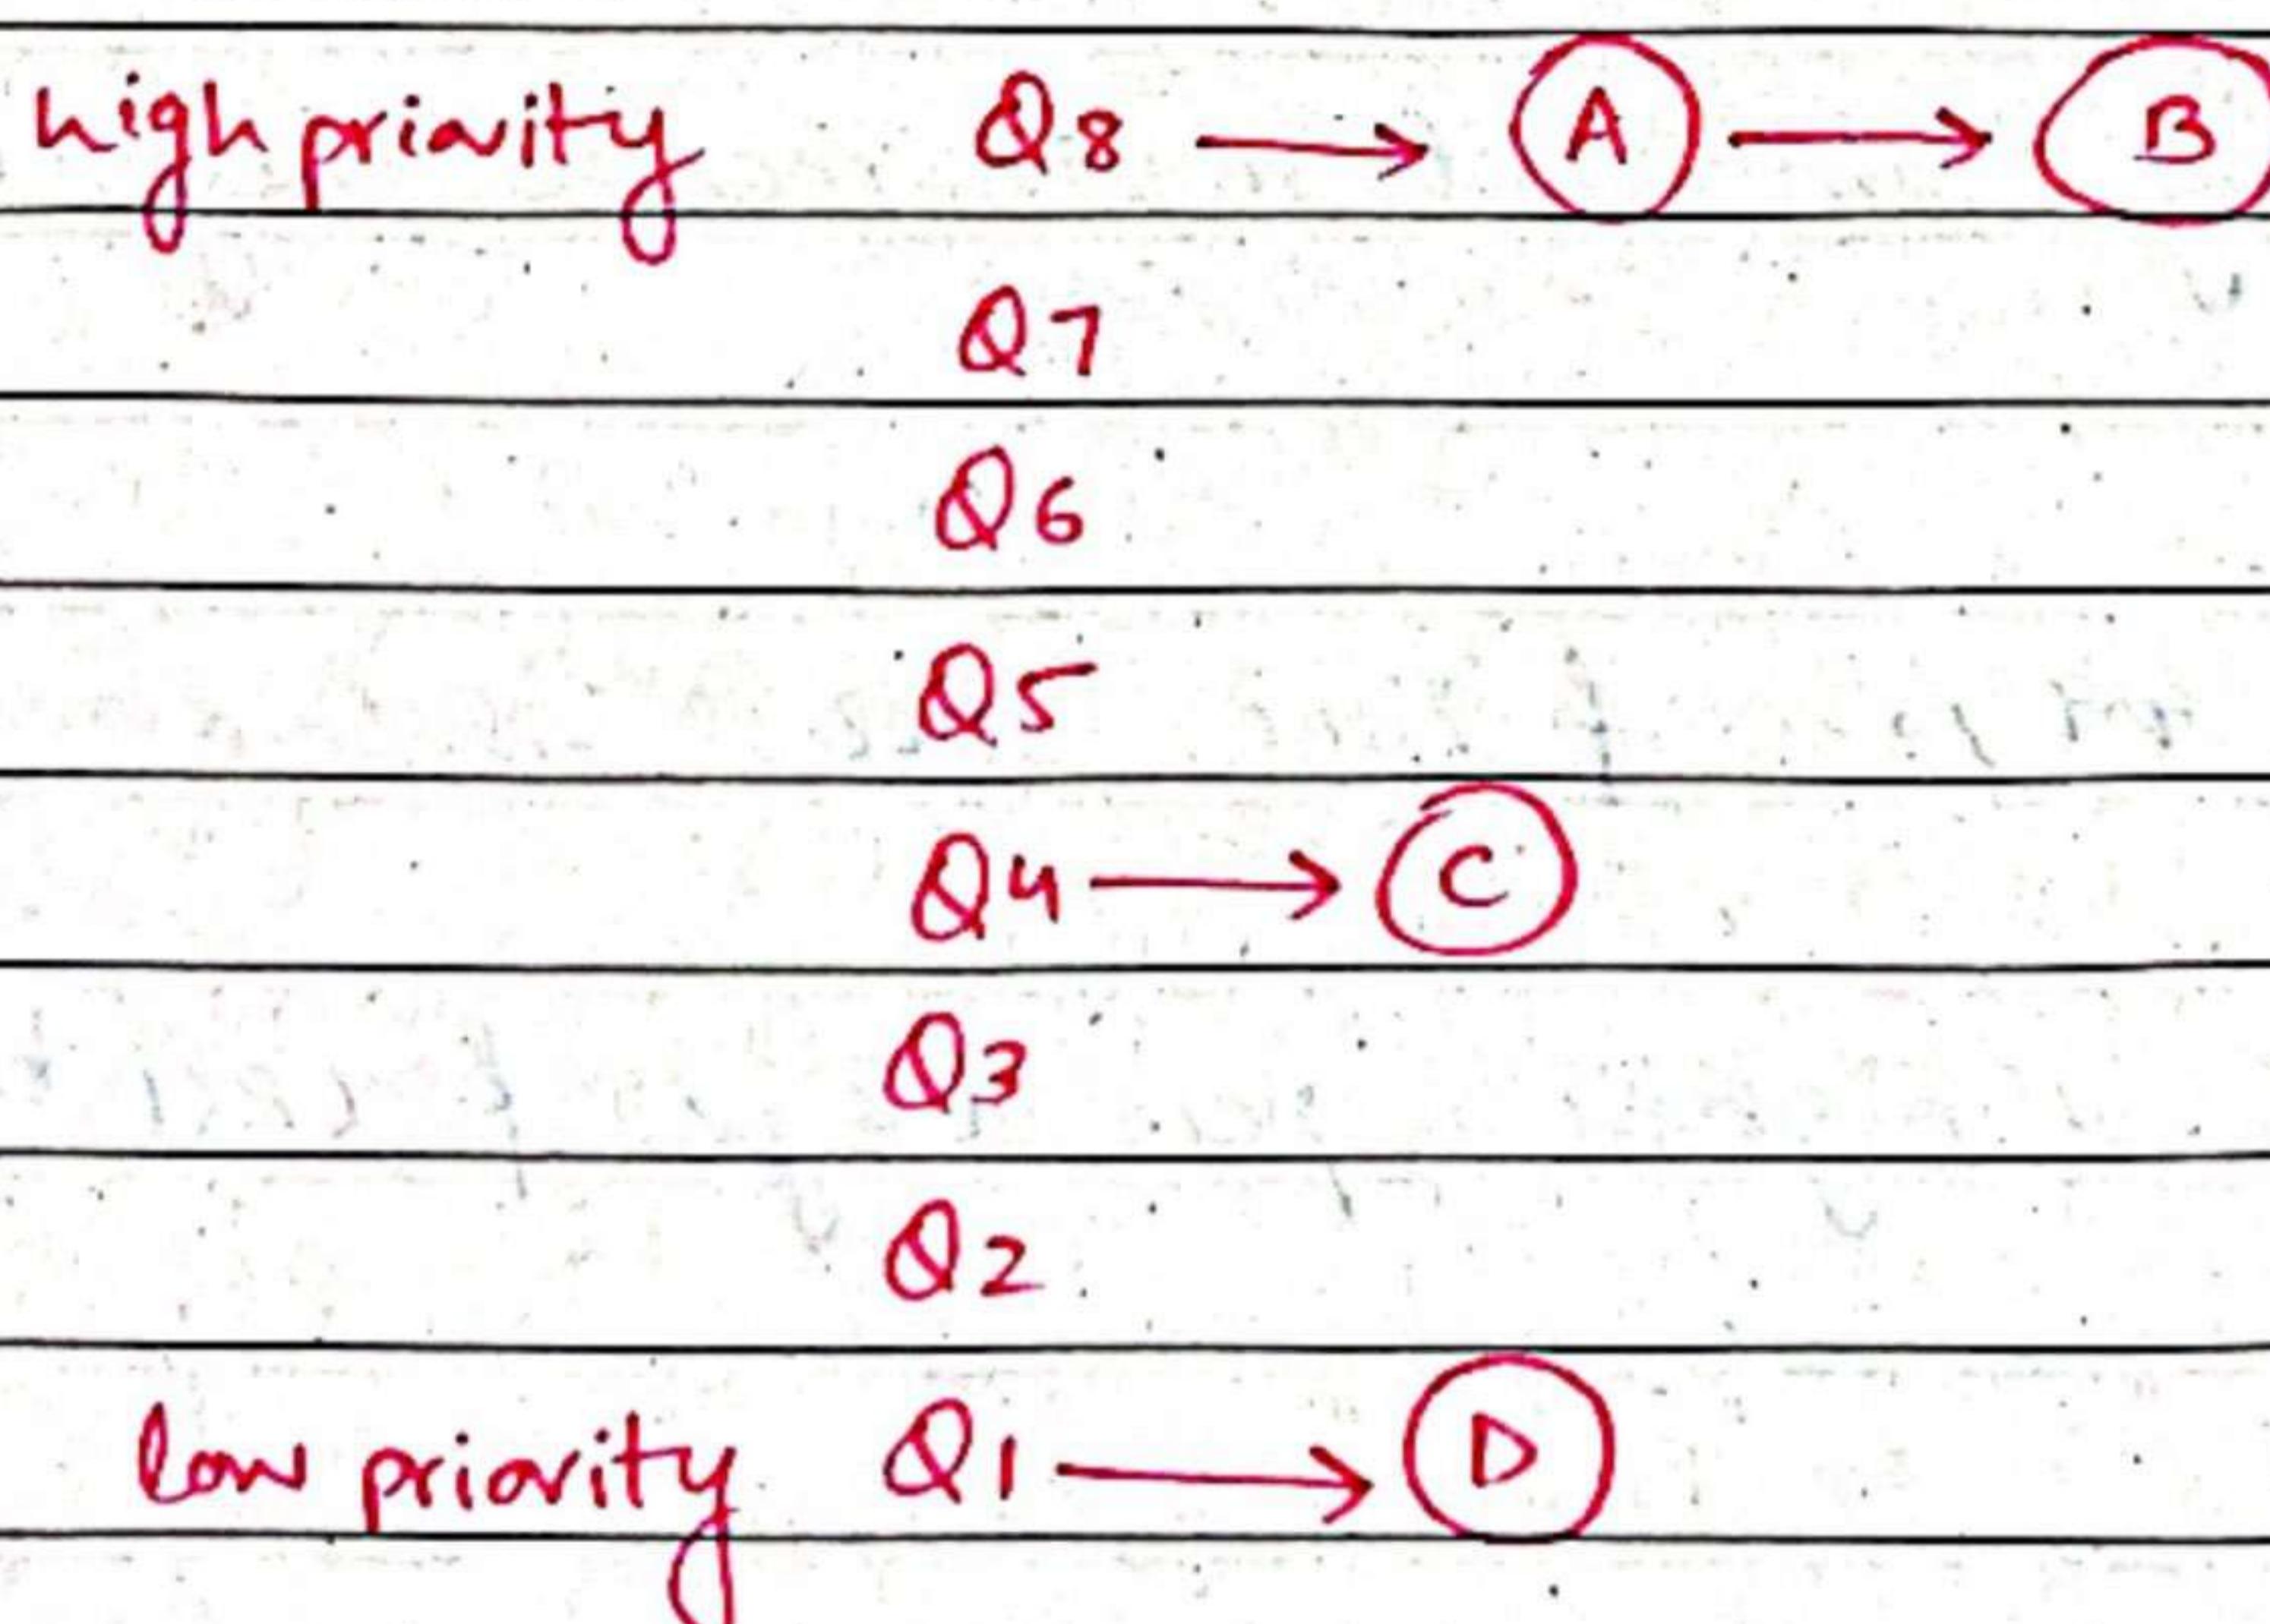
MLFQ : Basic Rules

- MLFQ has number of distinct distinct queues.
- each queue is assigned to different priority level.
- A job that is ready to run is on a single queue.
- a job on higher queue is chosen to run
- use round-robin scheduling among jobs in the same queue.

Rule 1: If Priority (A) > Priority (B), A runs (B doesn't)

Rule 2: If Priority (A) = Priority (B), A & B run in RR.

- MLFQ varies the priority of job based on its observed behaviour.



Date: _____

HOW TO CHANGE PRIORITY ?

→ MLFQ priority adjustment algorithm:

Rule 3: When a job enters the system, it is placed at highest priority.

Rule 4a: If a job uses up an entire time slice while running, its priority is reduced.

Rule 4b: If a job gives up CPU before time slice is up, it stays at same priority level.

In this manner, MLFQ approximates SJF.

PROBLEMS WITH MLFQ:

① Starvation

→ if there are "too many" interactive jobs in the system

→ long-running jobs will never receive any CPU time.

② Game Scheduler

→ After running 99% of time slice, issue an I/O operation

→ The job gains a higher percentage of CPU time.

Date: _____

- A program may change its behaviour time.
- CPU bound process → I/O bound process

Rule 5: After some time period S , move all the jobs in the system to topmost queue.

How to prevent gaming of our scheduler?

Solution

→ Rule 4: Once a job uses up its time allotment at given level, its priority is reduced.

Lower Priority, Longer Quanta.

→ The higher-priority queues → short time slices.

→ the low-priority queues → longer time slices.

SUMMARY

Rule ① If Priority (A) > Priority (B), A runs (B doesn't)

Rule ② If Priority (A) = Priority (B), A & B run in RR.

Rule ③ When a job enters the system, it is placed at the highest priority.

Rule ④ Once a job uses up its time allotment at given level (regardless of how many times it has given up CPU) its priority is reduced.

Rule ⑤ After some time period S , move all jobs in system to topmost queue.

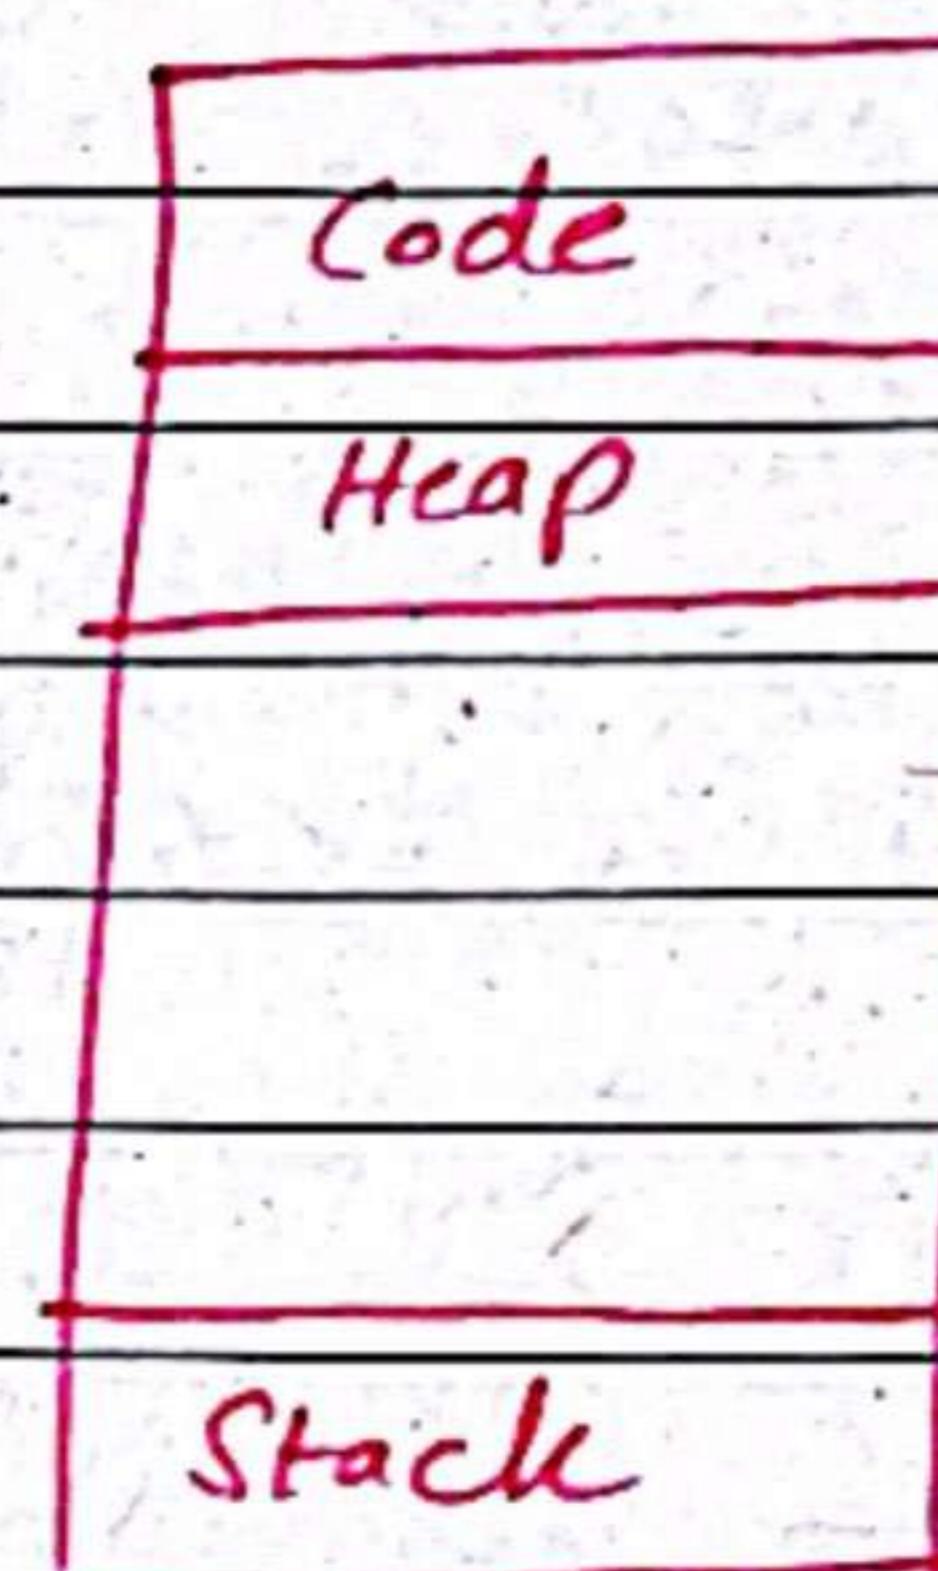
- Address space isolate process and provide illusion of dedicated memory.
 - The address space abstraction provides logical view of memory to a process
- global variables
Code → Data → Heap → Stack → local variables
Date: _____

CHAPTER 13: The Abstraction - Address Space.

→ An address space serves as an abstraction that provides each process with illusion of its own private memory

→ **MAIN COMPONENTS** of an Address Space:

- ① **Code section**: Contains the executable instructions of program
- ② **Stack**: Used for local variables, function parameters etc. arg[7] main function parameter
- ③ **Heap**: Provides memory for dynamic allocation (eg using malloc() in C)



HISTORICAL CONTEXT

- early computing systems used a single address space shared by all processes, leading to issues of memory protection and interference.
- Multiprogramming and time-sharing emerged as methods to allow multiple processes to run concurrently with isolated memory spaces.

IMPORTANCE OF ADDRESS SPACE

- Allows multiple programs to coexist in memory without interfering with each other.
- Enables OS to enforce boundaries for memory access, preventing a process from accidentally modifying another process's memory.

* OS is loaded at 0KB
in address space.

Date: _____

GOALS OF ADDRESS SPACES.

① **Transparency:** invisible to running program.

Each process operates as if it has its own contiguous block of memory, creating simplified programming model.

② **Protection:**

Processes are isolated from one another and memory access is restricted to prevent unauthorized interactions.

③ **Efficiency:**

Efficient memory management is essential to support multitasking, minimize overhead, and maximize resource utilization.

CHAPTER 14: Memory API

malloc()

void* malloc(size_t size).

* not a system call.
* found in #include <stdlib.h>

→ Allocate a memory region on the heap.

→ Argument

① size_t size: size of memory block (in bytes)

② size_t is an unsigned integer type.

→ Return

① Success: a void type pointer to the memory block allocated by malloc.

② Fail: a null pointer.