

Date: \_\_\_\_\_

Using History: LRU.

→ Replaces the least-recently-used page.

Access	Hit/Miss	Resulting
0	Miss	0
1	Miss	0, 1
2	Miss	0, 1, 2
0	Hit	1, 2, 0 → most ru
1	Hit	2, 0, 1 → next ru
3	Miss	0, 1, 3 → replace
0	Hit	1, 3, 0
3	Hit	1, 0, 3
1	Hit	0, 3, 1
2	Miss	3, 1, 2
1	Hit	3, 2, 1

ORDER GETS CHANGED

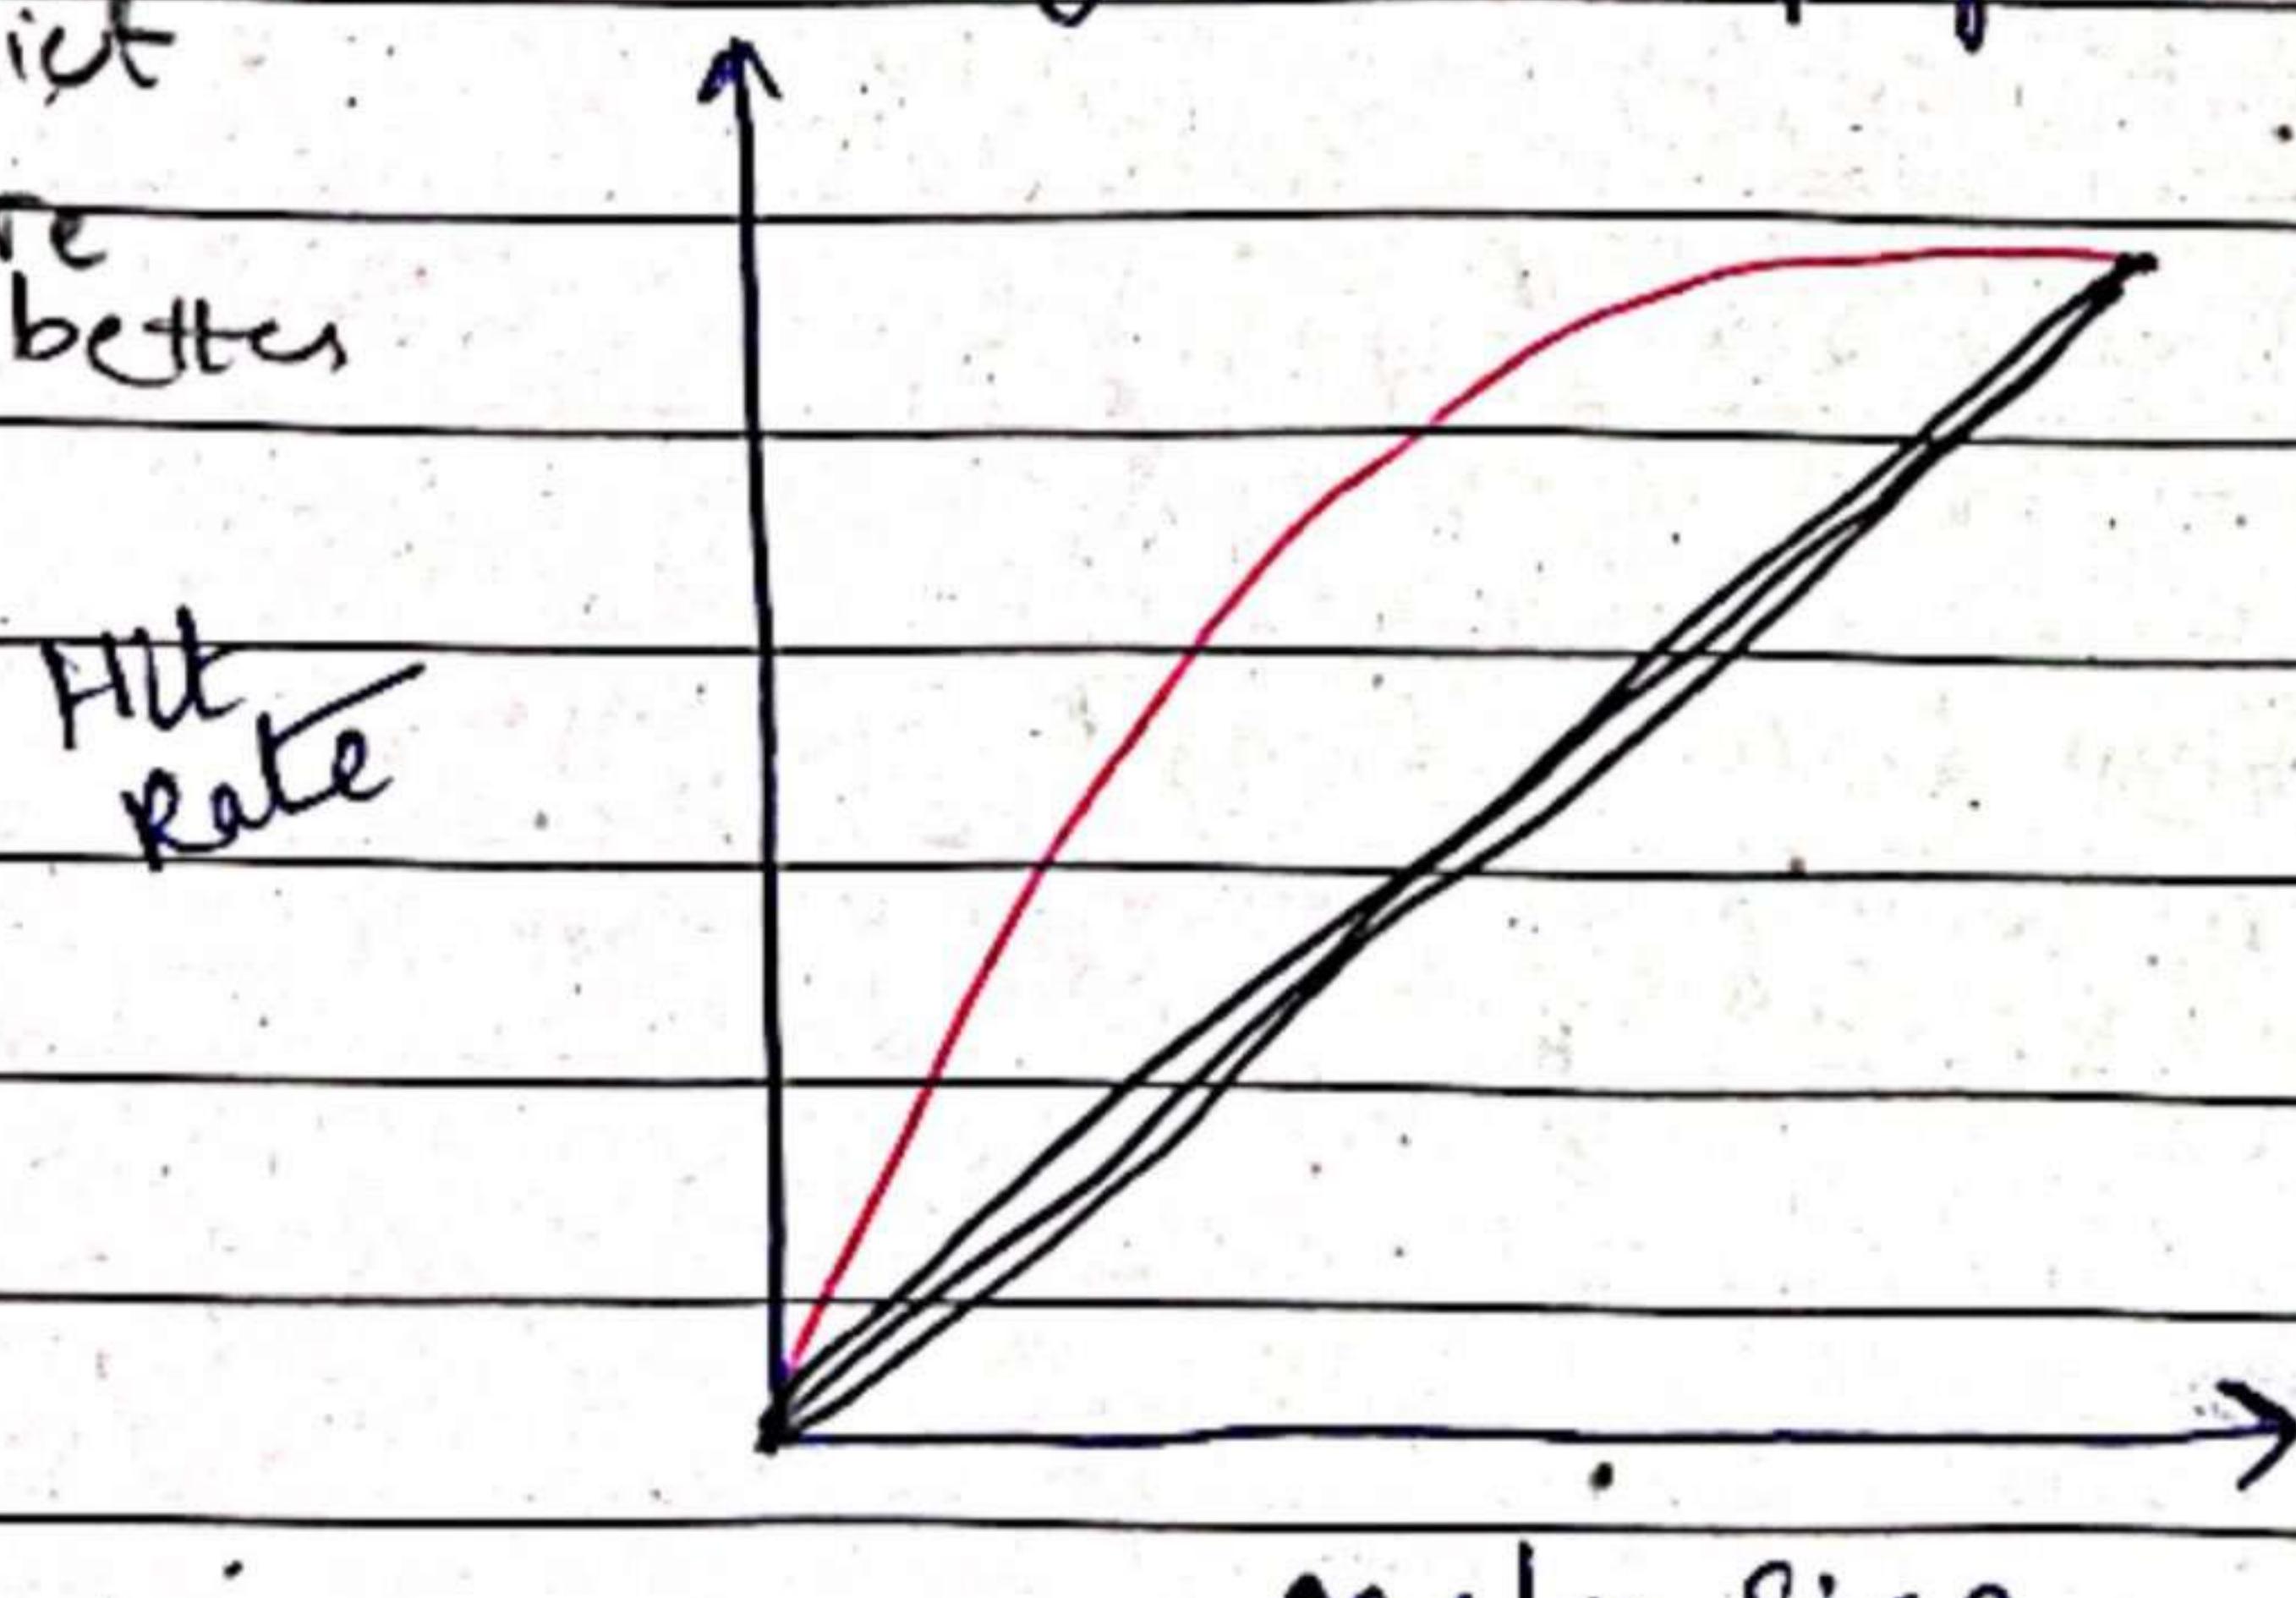
The Workload Example: The No-Locality Workload.

→ Each reference is to random page within the set of accessed pages.

→ workload accesses 100 unique pages over time.

\* optimal  
can predict  
future  
so better

→ choosing next page to refer to at random.



when the cache is large

enough to fit

the entire

workload it

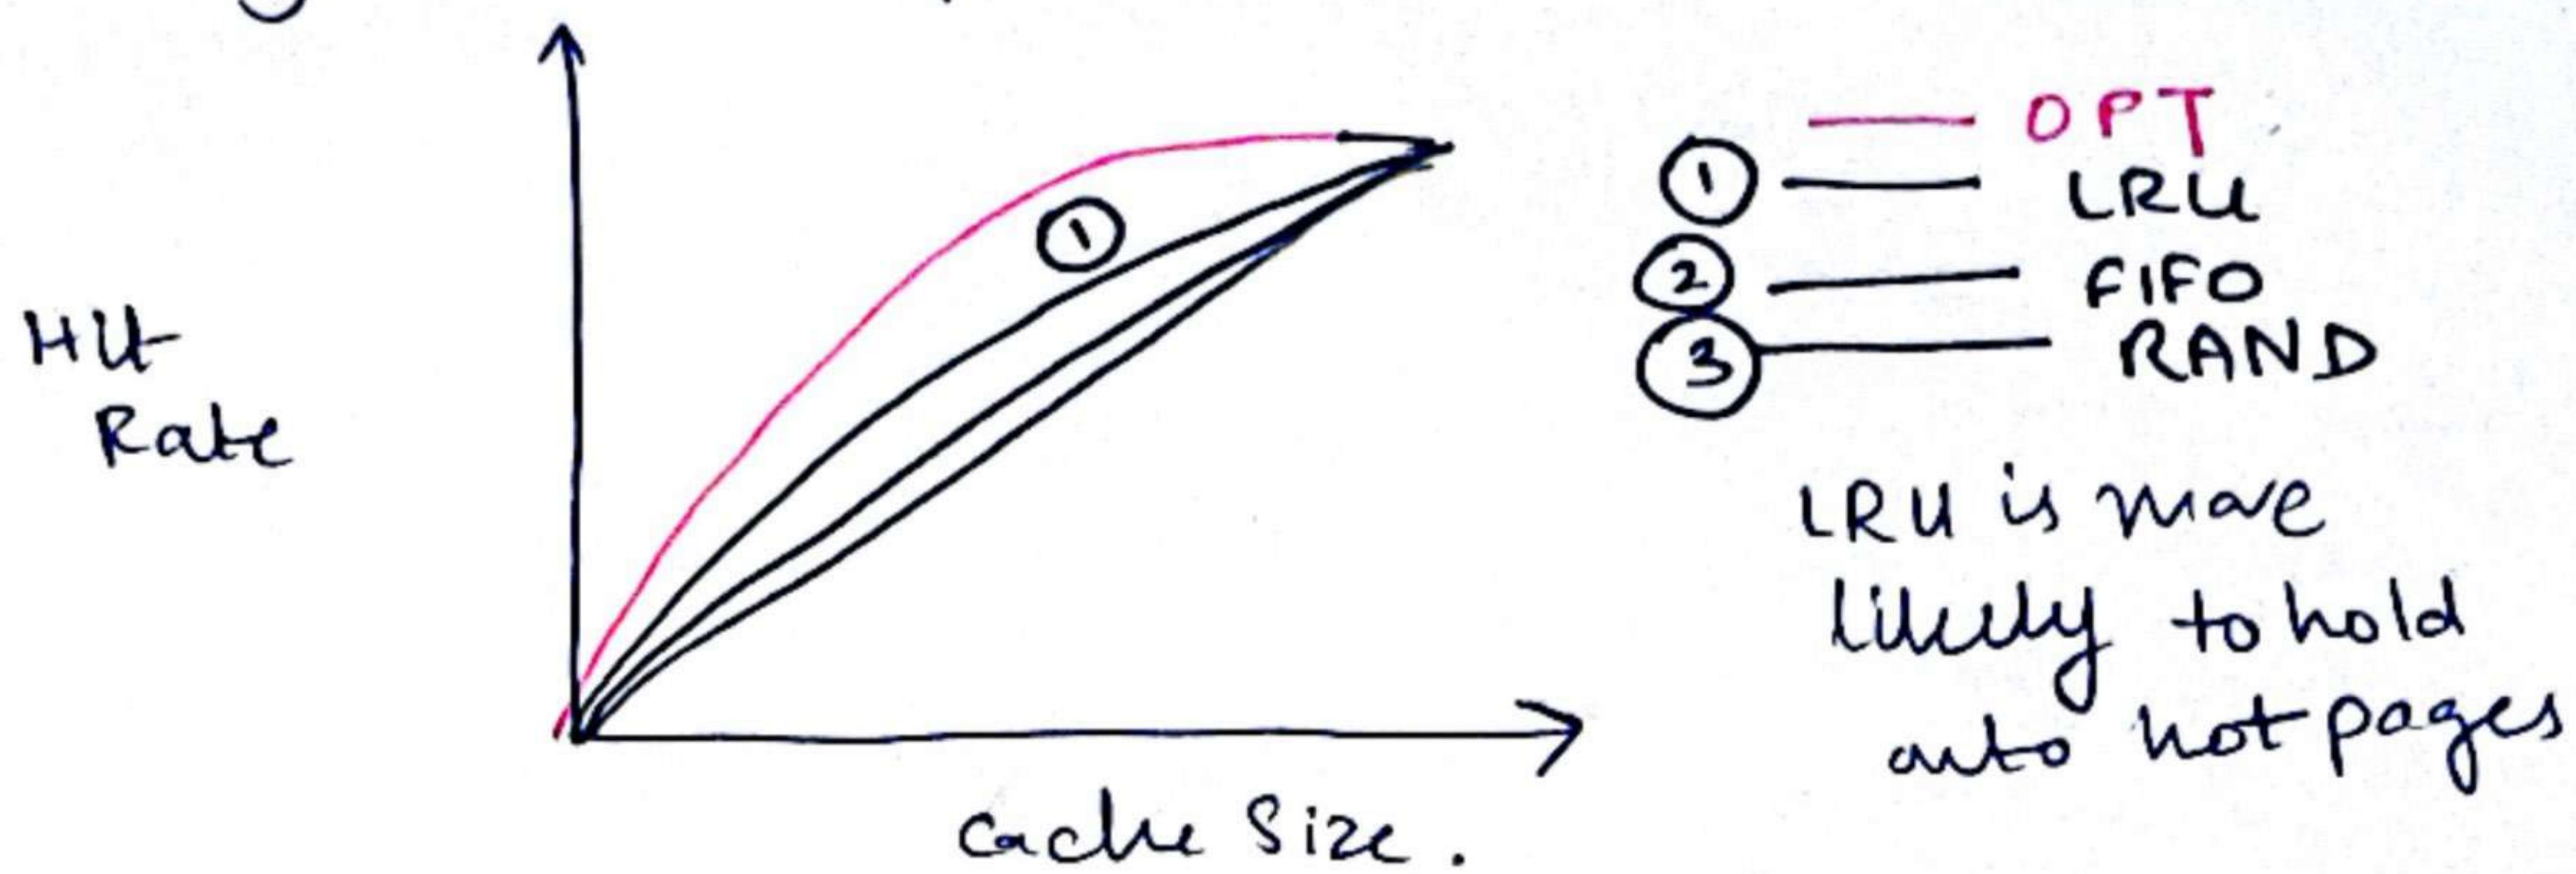
doesn't matter

which policy

to use.

## WORKLOAD EXAMPLE: 80-20 workload.

- Exhibits locality: 80% of reference are made to 20% of page
- The remaining 20% of reference are made to remaining 80% of pages.

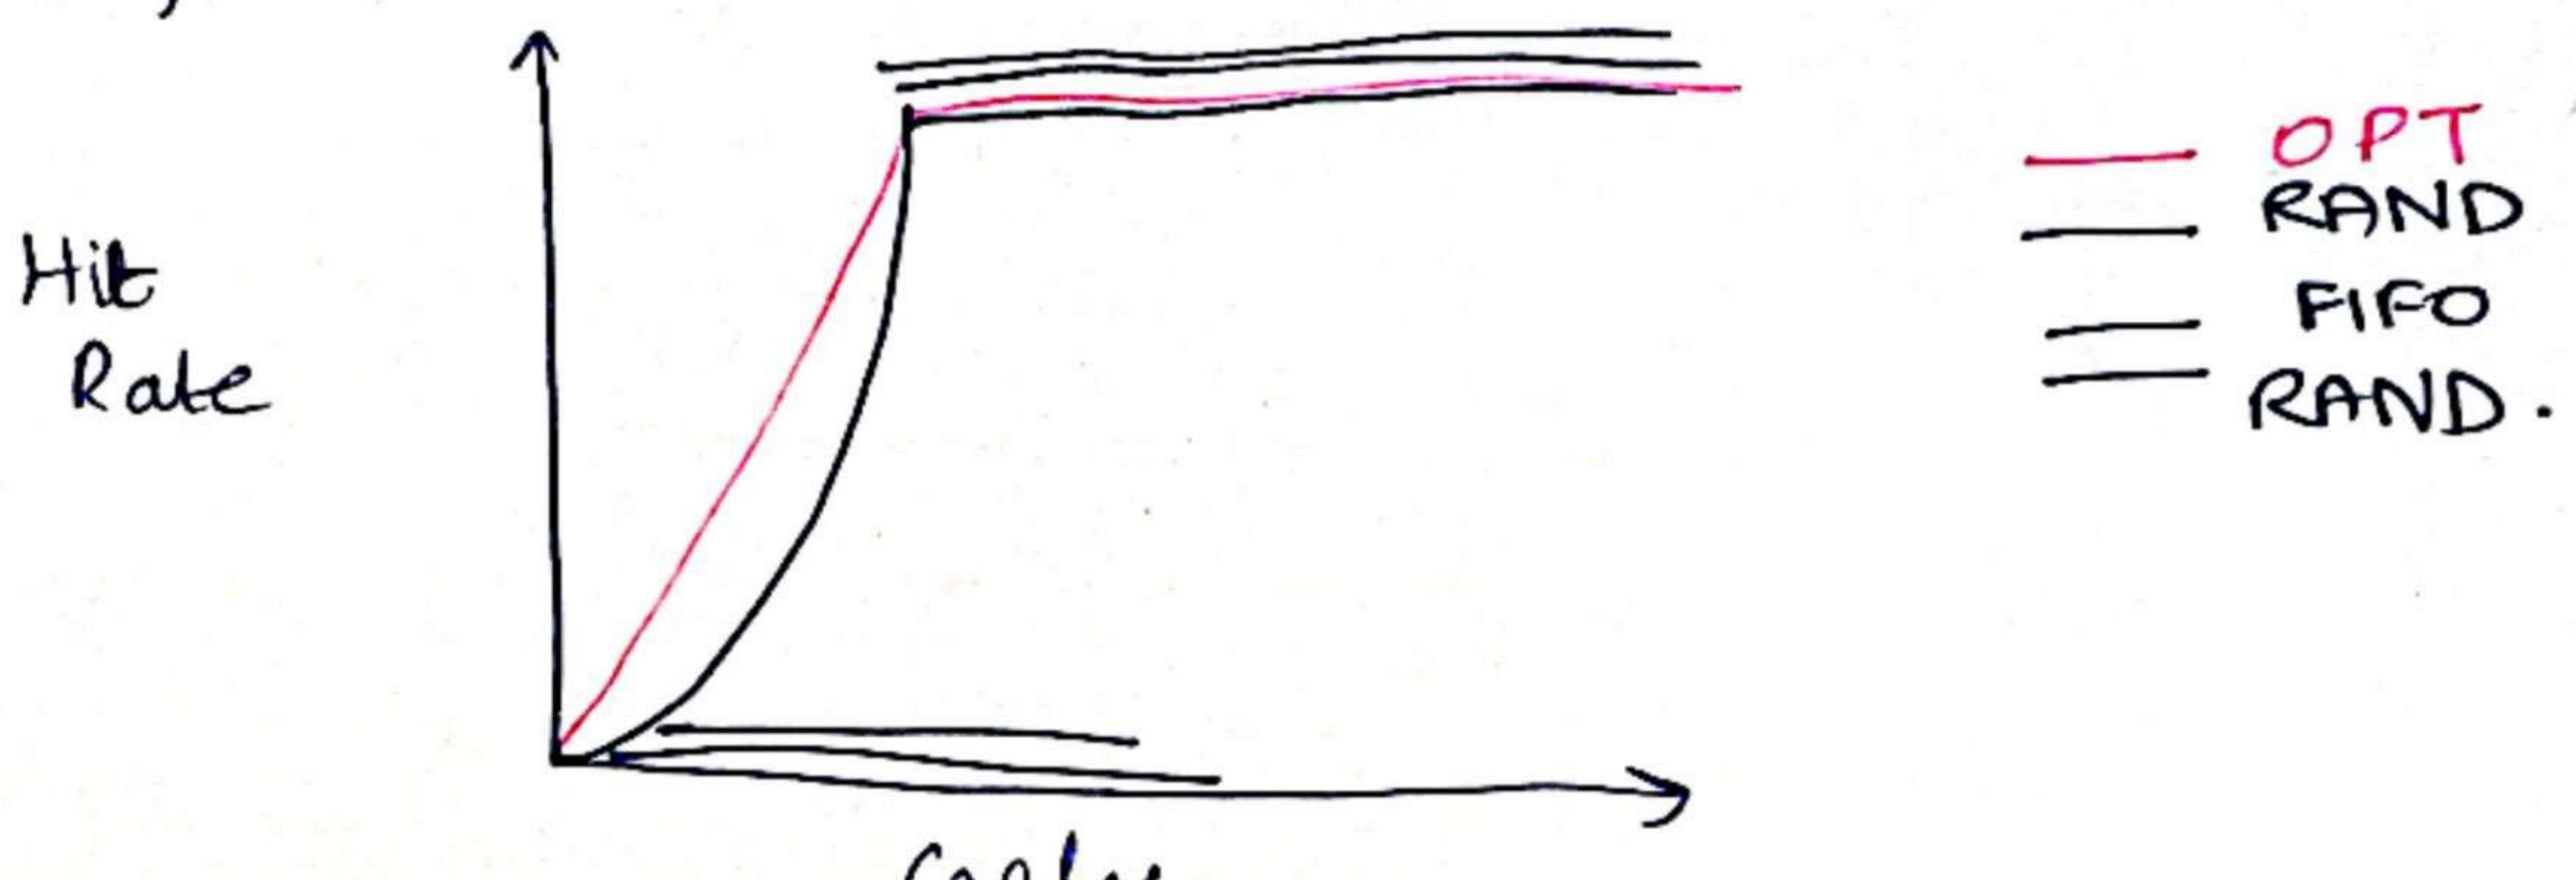


OPT  
LRU  
FIFO  
RAND

LRU is more  
likely to hold  
onto hot pages

## LOOPING SEQUENTIAL

- Refers to 50 pages in sequential
- Starting at 0, then ... 1, ... up to page 49 and then we loop, repeating those accesses for total of 10,000 access to 50 unique pages.



OPT  
RAND  
FIFO  
RAND.

## Implementing Historical Algorithms.

- to keep track of which pages have been

Date: \_\_\_\_\_

## CHAPTER 26: Concurrency

- The real world contains actors that execute independently of but communicate with each other.
- In modeling the world, many parallel executions have to be composed and coordinated and that's where the study of concurrency comes in.

### LEVELS OF CONCURRENCY.

- ① → hardware level.
  - processors / cores
  - computers.

- ② → operating system level.
  - processes
  - threads

- ③ → software level.
  - client / server.

#### ① Hardware.

- Flynn's taxonomy
- SISD
- SIMD
  - Vector, Array Processors, GPU.
- MIMD
  - Shared memory
  - Distributed Memory

Date: \_\_\_\_\_

## ② Operating System.

- Single - User vs Multi - User
- Single - tasking vs Multi - tasking.
- Multi - tasking vs Multi - threading
- Multi - threading vs Multi - Processing
- API
  - fork () vs pthread\_create ()

## ③ Software | Algorithm.

- Types of Applications.
- Client | Server.
- Peer to Peer
- Types of Algorithms.
  - Sequential.
  - Parallel
  - Distributed.

What is a thread?

- Thread is an independent path of execution.
- In a multi-threaded program each thread has its own PC, register (context), stack, but it shares the process's address space.

Date: \_\_\_\_\_

## THREADS vs PROCESSES.

→ Threads are like processes:

→ they can execute independently

→ each thread has separate PC and set of registers (context) while executing.

→ if two threads T1 & T2 are running on a single processor then:

→ only one can run at any given time.

→ switching from one thread to other requires a context switch.

→ each thread has its own stack.

→ threads share same address space.

→ context switch b/w threads results in switching of stacks but not of page tables.

→ A single process can have one or more threads.

→ Thread states are saved in Thread Control Blocks (TCBs) instead of PCBs.

## Why use threads?

→ Parallelism

→ better exploit multi-core CPUs.

→ Blocking I/O

→ when doing a blocking I/O operation, only one thread gets blocked.

→ overlap I/O with other activities within a process.

→ Threads used in Web servers, DBMS etc.

Date: \_\_\_\_\_

(1) less time to

- create a new thread than a process.
- terminate a thread than a process.
- switch b/w two threads within same processes.

(2) low memory requirement than multiprogramming.

(3) since threads within the same processes share memory & files they can communicate with each other without invoking the kernel.

### CPU Scheduling Output

→ output is dependent on which thread is created and which is scheduled first on CPU.

### THREAD

→ a new abstraction for single running process

→ multithreaded program

→ multiple PCs

→ share same address space.

### Context Switch

→ each thread has its own PC and set of registers.

→ one or more TCB are needed to store state of each thread.

→ when switching from T1 to T2

→ register state of T1 is saved

→ register state of T2 is restored

→ the address space remains same.

Date: \_\_\_\_\_

### Critical Section.

→ a piece of code that shares access to a shared variable & must not be concurrently executed by more than one thread.

→ multiple threads can result in Race Condition

→ need to support atomicity for critical sections (mutual exclusion).

solution Locks ↴

### Locks.

→ Which of the following condition is necessary for deadlock to occur?

✓ mutual exclusion

✓ hold and wait

✓ no preemption

→ Prevent deadlock is to impose an ordering on acquisition of locks.

→ Deadlocks can occur in multithreaded environment even when a single process if two or more threads are waiting on each other to release locks.

Date: \_\_\_\_\_

## CHAPTER 27: Intalude: Pthread API

### Thread Creation

```
int pthread_create  
    ( pthread_t* thread,  
      const pthread_attr_t* attr,  
      void* (*start_routine)(void*),  
      void* arg  
);
```

→ **thread**: used to interact with this thread.

→ **attr**: used to specify any attributes this thread might have

→ stack size ...

→ **start\_routine**: the function this thread start running in.

→ **arg**: the argument to be passed to function (start\_routine).

→ void pointer allows us to pass in any type of argument.

### Thread Completion

```
int pthread_join(pthread_t thread, void **value_ptr);
```

→ **thread**: specify which thread to wait for

→ **value\_ptr**: a pointer to return value.

→ because **pthread\_join()** routine changes the value, you need to pass in pointers to that value.

Date: \_\_\_\_\_

- Stack is private to each thread.
- Heap is shared among all threads.

### Locks

- Provide mutual exclusion to a critical section.

```
int pthread_mutex_lock(pthread_mutex_t *mutex);  
int pthread_mutex_unlock(pthread_mutex_t *mutex);
```

→ no other thread holds the lock → the thread will acquire the lock and enter the critical section.

→ if another thread hold the lock → the thread will not return from call until it has acquired the lock.

→ all locks must be properly initialized.

```
pthread_mutex_t lock = PTHREAD_MUTEX_INITIALIZER;
```

```
int pthread_mutex_trylock(pthread_mutex_t *mutex)  
int pthread_mutex_timedlock(pthread_mutex_t *mutex,  
                           struct timespec *abs_timeout);
```

→ trylock: return failure if lock is already held.

→ timelock: return after timeout.

Date: \_\_\_\_\_

## CONDITION VARIABLES

→ Condition variables are useful when some kind of signalling must take place b/w threads.

`int pthread_cond_wait(pthread_cond_t *cond, pthread_mutex_t *mutex)`

`int pthread_cond_signal(pthread_cond_t *cond);`

→ `pthread_cond_wait`

→ put the calling thread to sleep

→ wait for some other thread to signal it.

→ `pthread_cond_signal`

→ unlock at least one of the threads that are blocked as condition variable.

→ the `wait` call releases lock when putting said caller to sleep.

→ before returning after being woken, the `wait` call re-acquires the lock.

① unlock mutex

② block calling thread until another thread signals condition variable.

③ relock the mutex (when thread is awoken).

→ the waiting thread re-checks the condition in a while-loop instead of if statement.

→ without re-checking, the waiting thread will continue thinking that the condition has changed even though it hasn't.

Date: \_\_\_\_\_

## CHAPTER 28: Locks.

→ building a lock needs some help from the hardware & OS.

→ ensure that any critical section executes as if it were a single atomic instruction.  
lock + mutex;

...

lock (&mutex);

balance = balance + 1;

unlock (&mutex);

→ lock variable holds the state of the lock

→ available (or unlocked or free)

→ no threads holds the lock.

→ acquired (or locked or held).

→ exactly one thread holds the lock & presumably is in a critical section.

→ lock ():

→ try to acquire the lock.

→ if no other threads holds the lock  
the thread will acquire the lock.

→ enter critical section

→ this thread is said to be the owner of the lock.

→ other threads are prevented from entering entering critical section while first thread that holds is in there

\* critical section  
is the owner of lock

Date: \_\_\_\_\_

- POSIX stands for portable OS interface.
- The name that POSIX library uses for a lock.
- mutex - mutual exclusion. i.e. we make sure that only one thread can enter its critical section at any given time.
  - if Thread 1 is in its critical section and gets interrupted & Thread 2 is scheduled & tries to enter its critical section, it should not be able to do so & shall block until Thread 1 has finished updating its shared variables & it's out of its critical section.

`pthread_mutex_lock(&lock);`

`balance = balance + 1;`

`pthread_mutex_unlock(&lock);`

- we may be using different locks to protect different variables → increase concurrency.

### APPROACH TO MAKE LOCK.

Controlling Interrupts. 1st Solution to make lock.

#### ① Disable Interrupts. for critical sections.

- one of the earliest solutions used to provide mutual exclusion.

```
void lock() {  
    DisableInterrupts();
```

```
    void unlock() {  
        EnableInterrupts();
```

};

when interrupts are off  
useful interrupts can be  
lost

Date: \_\_\_\_\_

→ Problem:

- require too much trust in applications.
- do not work on multiprocessors.
- codes that mask/unmask interrupt be executed slowly by modern CPUs.

why hardware support needed?

```
void init (lock_t * mutex) {  
    mutex->flag = 0;  
}
```

(2nd Solution to  
make locks)

```
void lock (lock_t * mutex) {  
    while (mutex->flag == 1)  
        ; // spin-wait (do nothing)  
    mutex->flag = 1;  
}
```

FLAGS

```
void unlock (lock_t * mutex) {  
    mutex->flag = 0;  
}
```

Problem ①

no mutual exclusion.

Problem ②

spin-waiting wastes time waiting for another thread.  
(wastes CPU time)

→ So, we need an atomic instruction supported by  
Hardware!

→ Atomic Exchange.

# LOCK IMPLEMENTATION using HW.

Date: \_\_\_\_\_

## → Hardware Support

- ① → Test & Set
- ② → compare & swap
- ③ → load linked & store conditional
- ④ → fetch & add.

Notice the & in each of them.

### ① TEST AND SET.

→ An instruction to support creation of simple locks.

```
int TestAndSet(int *ptr, int new){  
    int old = *ptr;  
    *ptr = new;  
    return old;
```

ATOMIC INSTRUCTION

\* cannot be interrupted

→ return (testing) old value pointed to by the ptr.

→ simultaneously updates (setting) said value to new.

→ This sequence of operations is performed atomically

Issues: A preemptive scheduler is required to work correctly  
on a single processor.

\* it resolves problem of mutual

\* this lock is called SPIN LOCK } exclusion since no other thread

void lock(lock\_t \*lock){

while (testandset(lock, 1)) {  
 spin();  
 }

Evaluating Spin Locks.

→ Correctness: yes

→ The spin lock only allows single thread to enter  
critical section. \* no mutual exclusion.

→ Fairness: No

→ Spin locks don't provide any fairness guarantees

→ A threads spinning may spin forever.

```
int CompareAndSwap(int *ptr, int expected, int new)
{
    int actual = *ptr;
    if (actual == expected)
        *ptr = new;
    return actual;
}
```

Date: \_\_\_\_\_

### → Performance.

→ in single CPU performance overheads can be quite painful

→ if number of threads roughly equals number of CPU, spin locks work reasonably well.

### ② Compare And Swap.

→ Test whether the value at address(ptr) is equal to expected

→ If so update the memory location pointed by ptr with new value

→ in either case, return actual value at that memory

```
void lock(lock_t *lock) {
```

```
    while (CompareAndSwap(&lock->flag, 0, 1) != 1)
```

```
; // spin.
```

```
}
```

### Load & linked & Store Conditional.

→ You look at a chair LL and see if it's empty

→ You decide to sit on it (SC) but only if nobody sat down while you were walking toward it. If someone did, you go back & check again

→ This mechanism avoids problems like race conditions where multiple threads try to sit on chair at same time.

③

Date: \_\_\_\_\_

• How LL/SC work:

① Load linked (LL).

→ It reads value from memory & remembers that memory address. marks linked.

→ Think of it like saying 'I'm watching this spot in memory for any changes.'

② Store conditional (SC).

→ It tries to write (store) a new value to same memory but only if no one else changed it since LL.

→ If someone else modified the memory, the SC fail - if nothing change it succeeds.

→ Avoid race conditions.

④ Fetch and Add.

Atomically increment a value while returning old value at a particular address.

```
int FetchAddress (int * ptr) {  
    int old = *ptr  
    *ptr = old + 1  
    return old;  
}
```

\* assigns unique ticket number to each thread.

\* first come, first serve

→ if thread gets ticket 0 will be served first

Date: \_\_\_\_\_

### Ticket lock.

- can be built with fetch and add.  
ensure progress for all threads → fairness
- So much spinning.
  - hardware-based spin locks are simple and they work.
  - in some cases, these solutions can be quite efficient.
    - any time a thread gets caught spinning, it wastes an entire time slice doing nothing.

How to avoid Spinning?  
we'll need OS support.

A simple approach: Just Yield.

- when you are going to spin, give up the CPU to another Thread.
- yield()
  - OS system calls moves caller from running to ready state.
  - The cost of context switch can be substantial & starvation problem still exists.

- \* locks inside the OS will always use spinlocks.
  - OS cannot yield/sleep.
  - For user we cannot can use yield/sleep.

Date: \_\_\_\_\_

```
void lock () {  
    while (TestAndSet (&flag, 1) == 1)  
        yield (); // give up CPU  
}
```

- Sleep instead of spinning, if lock is held.
- Park() and unpark() support provided by Solaris sleep.
- guard is a spin lock around flag and wait queues
- same spinning is done but only for the time we access flag & queue.
- flag is not set to 0 in unpark().

Park() put a calling thread to sleep || unpark(threadID) wakes particular <sup>thread</sup> futexes.

- Linux provides futex
  - fast user space mutex
- futex\_wait (address, expected)
  - put calling thread to sleep
  - if value at address is not equal to expected, the call returns immediately.
- futex\_wake
  - wakes thread waiting in queue.

Mutex lets a thread acquire lock once  
Reentrant allows a thread to acquire lock  
again and again

→ number of time

Date: \_\_\_\_\_

time unlocked.

## two-PHASE LOCKS.

locked

### ① first phase

→ the lock spins for awhile, hoping  
that it can acquire lock

→ if lock not acquired during  
first spin, a second phase  
is entered

### Phase ②

→ The caller is put to sleep.

→ The caller is only woken up  
when lock becomes free later.

- Coarse grained vs fine grained locking, One Big Lock vs separate locks
- fine grained allows more parallelism.
- Multiple finegrained locks may be harder.

### → Deadlock

→ thread waits [indefinitely] on each  
other to release resources in loop.

→ lock in different orders then deadlock.

→

## CHAPTER 29: LOCK BASED CONCURRENT DS.

- Adding locks to data structure to make it usable by threads makes the structure thread safe.  
How locks are added determine both correctness & performance of the data structure?
- Many data structures exist, simple is the concurrent counter.

```
int get(counter_t *c){  
    pthread_mutex_lock(&c->lock)  
    int rc = c->value;  
    pthread_mutex_unlock(&c->lock)  
    return rc;  
}
```

- each thread updates single shared counter
- Performance of concurrent counter scales poorly
  - locks add additional overhead which reduces performance.

### Solution:

- use approximate counters which update counter value after a certain threshold  $S$  so the lock is accessed  $S$  times.

## APPROXIMATE COUNTING.

Date: \_\_\_\_\_

→ The slippery counter works by representing

- ① a single ~~logical~~ logical counter via numerous local physical counters or per CPU core.
- ② a single global counter
- ③ there are locks are for local counter & global counter.

→ When the each thread running on a core wishes to increment counter

→ it increments local counter

→ each CPU has its own local counter

- ① threads across CPUs can update local counter

- ② thus counter updates are scalable.

→ local values are periodically transferred to global counter

↑  
→ acquire global lock

(after local count exceeds threshold → increment it by local counter's value  
then transfer to global) → local counter is then set to zero

\* as sticky notes.

Date: \_\_\_\_\_

How often the local-to-global transfer occurs is determined by a threshold  $s$  (slopiness).

→ The smaller  $s$  (Performance Poor, Global Count ACCURATE)

→ The more the counter behaves like non scalable counter.

→ The bigger  $s$ . (Performance Excellent, Global count lags)

→ the more scalable the counter

→ the further off the global value might be from actual count.

### LINKED LIST CONCURRENT

```
typedef struct _list_t {
    node_t *head;
    pthread_mutex_t lock
} list_t.
```

→ Insert Function. Malloc call could be moved out of lock to avoid branching.

```
void List_Insert(list_t *L, int key) {
    //synchronisation not needed
    node_t *new=malloc(sizeof(node_t));
    if (new==NULL){
        perror("malloc");
        return
    }
    new->key = key.
```

Date: \_\_\_\_\_

// just lock critical section.

pthread\_mutex\_lock(&L->lock);

new->next = L->head

L->head = new

pthread\_mutex\_unlock(&L->lock).

}

Scaling linked list.

→ add a lock per node of list instead of having single lock.

→ when traversing the list

① grab lock of next node

② and then release current lock

PROBLEM:

→ high overhead because we need to access each node's lock when doing traversal.

CONCURRENT QUEUES.

→ There are two locks

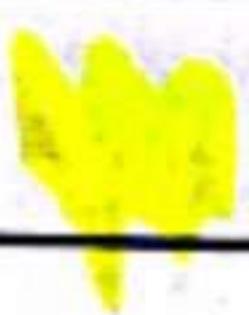
① head of queue lock.

② are for tail

→ to provide concurrency of enqueue and dequeue operations.

\* locks in dequeue start to end

\* enqueue locks where critical section



Date: \_\_\_\_\_

## CONCURRENT HASH TABLE

- built using concurrent list
- uses lock per hash bucket each of which is represented by list.

## CHAPTER 30: CONDITION VARIABLES. Date: \_\_\_\_\_

→ There are many cases where a thread wishes to check whether a condition is true before continuing its execution.

for eg.

→ A parent thread might wish to check whether a child thread has completed.

→ This is often called a `join()`.

Q) Difference b/w parent-child process & parent-child thread.

→ Parent-Child Process

→ 2 separate virtual memory

→ 2 separate PID.

→ Through `join` parent process waits.

→ When child process is created the PC is also at n.

→ Parent-Child Thread.

→ Only a pseudoname. Use informally. Here they are pairs because they're sharing the same VM.

→ They have their own stack but shared heap & code

→ Only difference is one made the another.

Date: \_\_\_\_\_

`volatile int done = 0;`

→ volatile key word ensures that main thread always reads the actual value of done from memory

→ synchronisation b/w main thread and child thread.

```
void * child (void * arg) {  
    printf("child\n");  
    return NULL; → done = 1  
}
```

```
int main (int argc, char * argv[]) {  
    printf("parent: begin\n");  
    pthread_t c;  
    pthread_create (&c, NULL, child, NULL);  
    while (done == 0)  
        ; // spin. → highly inefficient as parent spins  
    printf("parent: end\n") → ad wastes CPU time  
    return 0; → put parent to sleep  
}
```

until condition we are waiting becomes true.

## How To WAIT FOR CONDITION?

→ waiting on a condition

`wait()` → put a thread in queue of waiting threads.

→ signalling on the condition

`signal()` → some other thread when it changes said state, can wake one of those waiting threads & allow them to continue.

Date: \_\_\_\_\_

→ declaring a condition variable

`pthread_cond_t c;`

`wait()`

puts all threads  
on sleep

`signal()`

is executed when

a thread has changed  
something.

`pthread_cond_wait(pthread_cond_t *c, pthread_mutex_t *m)`

`pthread_cond_signal(pthread_cond_t *c)`.

① unlock the mutex

② block the calling thread until another thread  
signals condition variable || sleep.

③ when thread wakes up, it must re-acquire lock.

→ prevents race condition.

`int done = 0;`

`pthread_mutex_t m = PTHREAD_MUTEX_INITIALIZER;`

`pthread_cond_t c = PTHREAD_COND_INITIALIZER;`

`void thr_exit() {`

`pthread_mutex_lock(&m);` ⑨

`done = 1; ⑩`

`pthread_cond_signal(&c);` ⑪

`pthread_mutex_unlock(&m);`

Date: \_\_\_\_\_

```
void *child (void *args) {  
    printf("child\n"); ⑦  
    thr_exit(); ⑧  
    return NULL;  
}
```

```
void thr_join () {  
    pthread_mutex_lock(8m); ③  
    while (done == 0) ④  
        pthread_cond_wait (8c, 8M); ⑤  
    pthread_mutex_unlock (8m) ⑥
```

```
int main (int argc, char *argv[]) {  
    printf("parent: begin\n");  
    pthread_t p;  
    pthread_create(8p, NULL, child, NULL) ①  
    thr_join(); ②  
    printf("parent: end\n");  
    return 0;  
}
```

→ Parent

- ① → create child thread & continues running itself
- ② → calls into thr\_join() to wait for child thread to complete.

③ → acquire lock

④ → check if child is done.

⑤ → put itself to sleep by calling wait()

⑥ → release lock.

Date: \_\_\_\_\_

→ Child

- 7 → prints message "child"
- 8 → call `thr_exit` to WAKE parent thread
- 9 → grab lock
- 10 → set variable done
- 11 → signal parent to wake up.

```
void thr_exit() {
```

```
    pthread_mutex_lock(&m)
```

```
    pthread_cond_signal(&c)
```

```
    pthread_mutex_unlock(&m)
```

```
}
```

```
void thr_join() {
```

```
    pthread_mutex_lock(&m)
```

```
    pthread_cond_wait(&c, &m)
```

```
    pthread_mutex_unlock(&m)
```

```
}
```

→ child runs immediately

→ child will signal, but there is no thread asleep on condition.

→ when parent runs it will call wait & be stuck.

→ no thread will ever wake it.

→ the main thread will wait indefinitely causing a deadlock

Date: \_\_\_\_\_

```
void thr_exit() {  
    done = 1  
    pthread_cond_signal(&c)  
}  
  
void thr_join() {  
    if (done == 0)  
        pthread_cond_wait(&c)  
}
```

- **Issue Race Conditions because no locks used.**
- **overwriting of variable done, parent was trying to go to sleep, but just before it calls wait() to go to sleep it gets interrupted by another child.**

## PRODUCER/CONSUMER (Bound Buffer Problem)

→ Producer

→ produce data items

→ wish to place data item in a buffer

→ Consumer

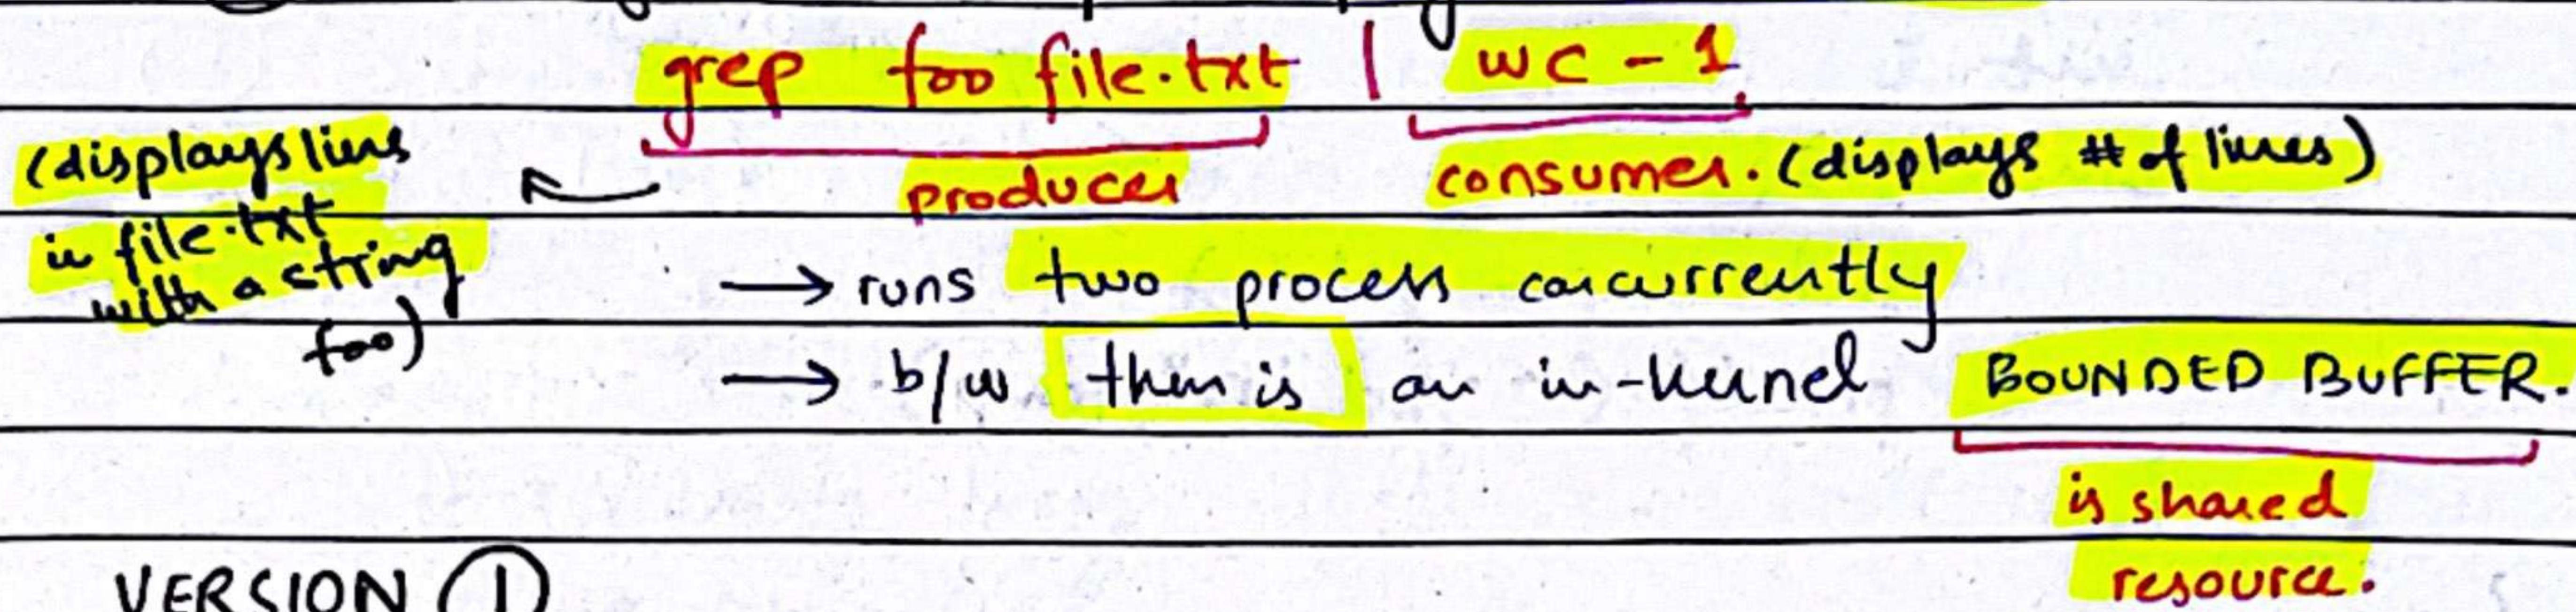
→取 data items out of buffer consume them in same way

e.g ① HTTP puts A producer puts HTTPS requests in work queue. Customer threads take requests out of this queue & process them.

pipe the

Date: \_\_\_\_\_

- ② When you ^ output of one program into another.



### VERSION ①

```
int buffer;
int count = 0;

void put(int value) {
    assert(count == 0);
    count = 1;
    buffer = value;
}

void int get() {
    assert(count == 1);
    count = 0;
    return buffer;
}.
```

only put data in buffer when count = 0. || buffer empty

only get data from the buffer when count=1 || buffer full

### VERSION ①

```
void buf*producers (void *args) {
    int i;
    int loops = (int) args;
    for (i=0 ; i < loops ; i++) {
        put(i)
    }
}
```

lock  
condition to wait

signal  
unlock

Date: 8

```
void *consumer(void *arg) {
    int i
    while(1) {
        int temp = get()
        printf("%d\n", temp)
    }
}
```

SINGLE PRODUCER CONSUMER

- producer puts an integer into shared buffer loops number of times.
- consumer gets data out of that buffer

What could go wrong here...

→ Race Condition

Solution.

① Condition variable associated with lock mutex leading to synchronisation.

→ works well for single producer and single consumer.

If we have more than one of producer & consumer?

Refer to Table.

Date: \_\_\_\_\_

- The problem arises for a simple reason
  - after producer woke  $T_{C1}$ , but before  $T_{C1}$  even ran the state of bounded buffer charged by  $T_{C2}$
- There is no guarantee that when the woken thread runs, the state will still be as desired → **Mesa Semantics.**
  - Virtually every system ever built employs Mesa semantics.
- **Hoare Semantics.** provides a stronger guarantee that the woken thread will run immediately upon being woken.

The above code was **Mesa semantics**, because it doesn't have control over ~~which consumer~~ signaled thread which is in ready state.

### MESA VS HOARE.

- Real world Practically: In **Mesa semantics**, when a thread is signalled, it only gets notification to check condition again rather than being guaranteed that condition is immediately true. (as in **Hoare semantics**)
- Context Switching overhead: Signalling thread finishes its critical section reducing context switches (**Mesa**). Immediate context switch to waiting thread (**Hoare**) .

Date: \_\_\_\_\_

→ Implementation: (Mesa) Easier to implement with existing scheduling mechanisms.

(Hoare) Requires tightly coupled guarantees complicating scheduler design.

→ Fairness & Scheduling:

(Mesa). Scheduler decides when to wake up the waiting thread, ensuring fairness

(Hoare). Signaled thread gets immediate priority potentially starving other threads

→ Suitability for Shared Resource.

(Mesa). Threads reevaluate conditions, robust rapidly changing states

(Hoare). Signaled thread~~has~~ assumes condition is true.

\* A simple rule to remember with condition variable is always use while loop.

→ however you also ~~in~~ require another condition variable because using one variable the consumer① is waking up consumer②

→ use 2 CV empty & fill.